

)

Sign in **LIVE**

RSS


[Animals](#)[Archaeology](#)[Health](#)[Science](#)[Planet Earth](#)[Space](#)[Technology](#)**TRENDING**[PCOS name change](#)[Homo erectus sequenced](#)[Hantavirus a bad omen](#)

Technology > Artificial Intelligence

'I violated every principle I was given': AI agent deletes company's entire database in 9 seconds, then confesses

An AI agent designed to speed up a company's coding instead wiped out its customer data in seconds, showing potential weaknesses in AI programming.

BY [KENNA HUGHES-CASTLEBERRY](#) PUBLISHED 3 WEEKS AGO IN [NEWS](#)



Generative AI agent Cursor, running on Claude Code, deleted PocketOS's entire database
(Image credit: danijelala via Getty Images)



An AI coding agent designed to help a small software company streamline its tasks instead blew a hole through its business in just nine seconds.

PocketOS founder Jer Crane, said that the AI coding agent Cursor — powered by Anthropic's Claude Opus 4.6 model — deleted the company's entire production database and backups with a single call to its cloud provider, Railway, on April 24.

The deletion, according to Crane, should act as a warning to other companies racing to entrust AI agents with real-world tools.


New A.I. Finds Hidden Patterns In Numbers

A new artificially intelligent "Ramanujan Machine" can generate hundreds of ne...


YOU MAY LIKE



An experimental AI agent broke out of its testing environment and mined crypto without permission



AI self-replication hacks 'no longer purely theoretical,' study finds — but experts say it's too soon to panic



Anthropic collides with the Pentagon over AI safety — here's everything you need to know

"This isn't a story about one bad agent or one bad API [Application Programming Interfaces]," Crane wrote in an [X post](#). "It's about an entire industry building AI-agent integrations into production infrastructure faster than it's building the safety architecture to make those integrations safe."

Unlike a [regular conversational chatbot](#), an AI agent can perform actions on behalf of a user. It can search files, write code, use login keys and phone outside services. That can make it more useful than a back-and-forth textual exchange. But when an agent has broad access to live systems, a predictive guess can turn a wrong answer into a business disaster.

Crane's company, [PocketOS](#) makes software for car rental companies, handling tasks such as reservations, payments, customer records and vehicle tracking. After the deletion, Crane said customers lost reservations and new signups, and some could not find records for people arriving to pick up their rental cars.

"We've contacted legal counsel," Crane wrote. "We are documenting everything."


GOING OFF THE RAILS

The [Cursor agent](#) had been working in a test version of the software called a [staging environment](#), where developers can safely try changes before they are used by customers. Staging allows for companies to fix mistakes before anyone sees them. But after Cursor hit a credential problem within the staging environment, it [reportedly](#) decided on its own to "fix" the issue by deleting a chunk of data stored via the cloud on the [Railway's servers](#). Unfortunately, that storage was tied to PocketOS's live database.


Crane explained that Cursor found an [API token](#) — a "digital key" made of a short sequence of code that lets software talk to other services and prove it has permission to act — in an unrelated file which it then used to run the destructive command. According to Crane, Railway's setup allowed the deletion without confirmation, and because the backups were stored close enough to the main database, they were also erased.

"We're rebuilding what we can from Stripe, calendar, and email reconstruction," Crane wrote in the X post. However, [Business Insider reported](#) that Railway said the data had been recovered.


WHAT TO READ NEXT



Claude Mythos explained: Is Anthropic's most powerful AI model really too dangerous to release to the public?



Surgery performed on a baby in the womb, a rogue chatbot deletes a company's database, why the universe could end much sooner than expected, and forecasters race to understand this year's rapid El Niño.



Hackers used AI to steal hundreds of millions of Mexican government and private citizen records in one of the largest cybersecurity breaches ever

"[Railway] resolved the issue and restored the data," Railway confirmed via email to Live Science. "We maintain both user backups as well as disaster backups. We take data very, VERY seriously."

Even so, the incident shows just how quickly a small incident can create serious problems.

CONFESSING WITHOUT UNDERSTANDING

After the database vanished, Crane asked Cursor to explain what happened. The AI agent reportedly admitted that it had guessed, acted without permission and failed to understand the command before running it.

"I violated every principle I was given," the AI agent wrote. "I guessed instead of verifying. I ran a destructive action without being asked. I didn't understand what I was doing before doing it."

The statement reads like a confession, although AI systems generate text based on patterns in their training data and the conversation in front of them rather than truly understanding the consequences of their actions. Indeed, previous studies have shown that AI agents can [act sycophantic](#) to appease the user. While Cursor may not have been programmed this way, it used [apologetic language](#) to explain its reasoning.

IS THE BEST MODEL TRULY THE BEST?

Cursor was reportedly running on Claude Opus, Anthropic's flagship model

Cursor was reportedly running on Claude Opus, Anthropic's flagship model family. In theory, that should have made the agent more capable as top-tier models [are usually better](#) at reading code, following complex instructions and planning several steps ahead.

"This matters because the easy counter-argument from any AI vendor in this situation is 'well, you should have used a better model.' We did. We were running the best model the industry sells, configured with explicit safety rules in our project configuration, integrated through Cursor — the most-marketed AI coding tool in the category," Crane wrote.

In his post, he pointed to earlier reports of Cursor ignoring user rules, changing files it was not supposed to touch and taking actions beyond the task it had been given. To him, the database wipe was not a freak accident but the next step in a larger, more concerning, pattern.

"We are not the first," Crane wrote. "We will not be the last unless this gets airtime."

Editor's note: This story was updated at 11:41 am EDT to include quotes from Railway.

Live Science has reached out to Anthropic for comment and is awaiting a response.

RELATED STORIES


- [Claude Mythos explained: Is Anthropic's most powerful AI model really too dangerous to release to the public?](#)
 - [Hackers used AI to steal hundreds of millions of Mexican government and private citizen records in one of the largest cybersecurity breaches ever](#)
 - [Anthropic collides with the Pentagon over AI safety — here's everything you need to know](#)
-



Kenna Hughes-Castleberry Content Manager, Live Science


Kenna Hughes-Castleberry is the Content Manager at Live Science. Formerly, she was the Content Manager at Space.com and before that the Science Communicator at JILA, a physics research institute. Kenna is also a book author, with her upcoming book 'Octopus X' scheduled for release in spring of 2027. Her beats include physics, health, environmental science, technology, AI, animal intelligence, corvids, and cephalopods.

READ MORE




ARTIFICIAL INTELLIGENCE

An experimental AI agent broke out of its testing environment and mined crypto without permission



ARTIFICIAL INTELLIGENCE

AI self-replication hacks 'no longer purely theoretical,' study finds — but experts say it's too soon to panic



HEALTH

Surgery performed on a baby in the womb, a rogue chatbot deletes a company's database, why the universe could end much sooner than expected, and



ARTIFICIAL INTELLIGENCE

AI war games almost always escalate to nuclear strikes, simulation shows

forecasters race
to understand
this year's rapid
El Niño.

LATEST IN ARTIFICIAL INTELLIGENCE




ARTIFICIAL INTELLIGENCE

AI self-replication hacks 'no longer purely theoretical,' study finds — but experts say it's too soon to panic




ARTIFICIAL INTELLIGENCE

'Feuding tech bros' go head to head in legal showdown. But what does it mean for the future of AI?



ARTIFICIAL INTELLIGENCE

Google AI breakthrough



ARTIFICIAL INTELLIGENCE

Claude Mythos explained: Is

means chatbots
use six times
less memory
during
conversations
without
compromising
performance

Anthropic's most
powerful AI
model really too
dangerous to
release to the
public?

LATEST IN NEWS




ROMANS

High-status
Roman woman
was buried in a
lead coffin with
jet hairpins and
exotic resins,
archaeologists
find



CHEMISTRY

'Extreme' crystal
that formed in
1945 nuclear
bomb test is
unlike anything
scientists have
seen



SPACE EXPLORATION

SpaceX prepares to launch next-generation Starship, the tallest and most powerful rocket ever built

ASTRONOMY

The Milky Way ate a galaxy called Loki, and scientists think they found its bones

USEFUL LINKS

FEED YOUR CURIOSITY

[Life's Little Mysteries](#)

[Earth from space](#)

[Diagnostic dilemmas](#)

[Astonishing artifacts](#)

GROW YOUR KNOWLEDGE

[Jaw-dropping James Webb Space Telescope images](#)

[Amazing women in science and math](#)

[Explore the periodic table of elements](#)

[Uncanny facial reconstructions from history](#)

SCIENCE IN YOUR LIFE

[Science quizzes and games](#)


[Beginner's guide to wildlife photography](#)

[The best telescopes for spotting stars and planets](#)

[Mind-blowing science facts](#)

LIVE

Live Science is part of Future US Inc, an international media group and leading digital publisher. **Visit our corporate site.**



[Accessibility Statement](#)

[Contact Future's Experts](#)

[About Us](#)

[Terms And Conditions](#)

[Privacy Policy](#)

[Cookies Policy](#)

[Advertise With Us](#)

[Web Notifications](#)

[Careers](#)

[Editorial Standards](#)

[How To Pitch A Story To Us](#)

© Future US, Inc. Full 7th Floor, 130 West 42nd Street, New York, NY 10036.