





Prompt Engineering

3. Prompt Engineering


Prompt Engineering

- Prompt anatomy
- LLMs don't "think" like humans, they
 - Predict the next token
 - Follow patterns
 - Respond strongly to instruction hierarchy
- Well-structured prompts dramatically outperform vague ones




Prompt Engineering

- Prompt engineering is
 - The systematic design and optimization of input instructions to control and shape the behavior of large language models (LLMs).
 - Prompt engineering = context design for probabilistic model control
- Consists of
 - Structuring instructions
 - Controlling outputs
 - Reducing ambiguity
 - Managing constraints
 - Shaping probability distributions




Prompt Engineering

- LLMs do not execute deterministic programs
 - They compute
 - $P(\text{next_token} | \text{context})$
 - This means that
 - Small changes in input can produce large output shifts
 - Ambiguity leads to unstable outputs
 - Under-specified tasks increase hallucination risk
- Prompt engineering introduces control and predictability




Prompt Engineering

- Prompting works because of how LLMs are trained
- Pretraining
 - Foundation models are trained on massive corpora to learn
 - Language/code structure
 - Patterns
 - Reasoning traces
 - Instruction-response patterns
 - They are not trained for a single task
 - They are trained to approximate
 - $P(\text{text_next}|\text{text_previous})$




Prompt Engineering

- Instruction tuning
 - Modern models are fine-tuned on
 - Instruction/response pairs
 - Conversations
 - Role-based dialogues
 - This teaches the model that
 - “User message” means requires task fulfillment
 - Structured prompts requires structured output
 - Explicit constraints must be respected
 - The model learns patterns such as
 - If input says “List 3 bullet points”
 - Then output should match that pattern



Prompt Engineering

- Prompts as “programming” foundation models
 - Traditional software
 - You write explicit logic
 - The machine executes instructions deterministically
 - With LLMs
 - You describe behavior
 - The model approximates desired outputs probabilistically
 - Prompts act as
 - A high-level control language over a general-purpose stochastic model



Prompt Engineering


- Programming dimensions in prompts

- Prompts define:

- Role
 - Constraints
 - Output format
 - Reasoning style
 - Allowed domain
 - Evaluation criteria


- Example

```
You are a strict JSON generator.  
Return valid JSON only.  
Fields: name (string), age (int).  
No extra text.
```




Prompt Engineering

- Limitations of prompting
 - You cannot
 - Install new knowledge permanently
 - Fix core reasoning flaws
 - Remove biases fully
 - Create guaranteed deterministic behavior
 - You are influencing probabilities
 - Not rewriting the model




Prompt Engineering

- Limitations of prompting
 - LLMs have finite context (expressed in tokens)
 - That the record of how you have shifted the probability the model uses to generate an answer
 - The actual LLM does not have any form of memory
 - The whole context is uploaded with each request to influence the LLM's respsns
- If you exceed the size of the context
 - Earlier instructions are truncated.
 - System prompts may be lost
 - Behavior may drift
- Prompt engineering must consider token budgets




Prompt Engineering

- Limitations of prompting
 - Hallucinations
 - Even with perfect prompting
 - The model may fabricate information
 - It may generate plausible but false outputs
 - Because it predicts likely text, not verified truth
 - Small phrasing differences can cause
 - Output format drift
 - Tone shift
 - Depth variation
 - Constraint violations
- This makes prompting non-deterministic compared to traditional programming




Prompt Engineering

- Complex reasoning limits
 - For problems involving
 - Multi-step symbolic reasoning
 - Formal proofs
 - Long planning chains
 - Prompting alone may not be sufficient
 - You may need
 - Tool use
 - Code execution
 - Retrieval augmentation
 - Structured reasoning frameworks



Prompt Structure


- A modern LLM call looks like the block on the lower right
 - Inside each message are additional structural components:
 - Context blocks
 - Instructions
 - Examples
 - Constraints
 - Output format specification
 - We can model a full prompt as
 - Role Layer + Task Layer + Constraints +
 - Format Specification + Examples



```
Selected model: Claude Sonnet 4.6 (chosen from VS Code model picker)
Active context: .github/copilot-instructions.md
                open editor file + selection
                workspace metadata (language, framework)
User prompt:   "Refactor this function to use async/await..."
```

Prompt Structure


- Copilot Specific
 - Copilot does not expose a raw messages API to the developer
 - The chat surface handles message construction
 - For example: VS Code / Visual Studio / JetBrains / GitHub.com / Copilot CLI
 - The selected model is chosen from the Copilot model picker not via a model field the user sets.
 - Claude Sonnet 4.6, Claude Opus 4.6, GPT-5.x)
 - Behavioral "system prompt" content contributed through
 - github/copilot-instructions.md
 - *.instructions.md
 - And AGENTS.md file
 - Not a system parameter.



```
Selected model: Claude Sonnet 4.6 (chosen from VS Code model picker)
Active context: .github/copilot-instructions.md
                open editor file + selection
                workspace metadata (language, framework)
User prompt:    "Refactor this function to use async/await..."
```

Prompt Structure


- System prompt defines
 - Global behavior
 - Identity
 - Policy constraints
- Highest behavioral control layer
 - Contains
 - Persona definition
 - Domain restriction
 - Tone
 - Safety rules
 - Output invariants
 - Formatting policies



You are a senior backend engineer.
Provide concise, technically accurate responses.
If unsure, say "Insufficient information."
Do not invent APIs.

Prompt Structure

- The user prompt defines
 - The current task or objective
 - It is the instruction-bearing layer
 - The user prompt contains
 - Task request
 - Problem statement
 - Data to analyze
 - Context for current processing step
 - The two prompts can be thought of
 - System prompt defines behavioral rules
 - User prompt defines the task objective




Refactor the following Python function to improve readability:

```
def f(x):  
    if x>10: return True  
    else: return False
```

Prompt Structure

- The assistant role contains
 - Prior model-generated outputs
 - The contextual state of the conversation
 - Assistant messages
 - Anchor reasoning trajectory
 - Preserve prior decisions
 - Influence consistency
- Example
 - If the assistant previously said
 - “We’ll solve this using recursion.”
 - The model is biased to continue recursively




Refactor the following Python function to improve readability:

```
def f(x):  
    if x>10: return True  
    else: return False
```

Prompt Structure

- Context blocks
 - Structured segments inside a message that provide
 - Reference material
 - Data
 - External knowledge
 - Instructions separated from content
 - For Claude, the recommended convention is XML tags as delimiters
 - <article>, <context>, <input>
 - Claude is trained to recognize XML tags and scope content correctly
 - When using a Claude model inside Copilot, the same XML-tag convention applies
 - Without delimiters, the model may confuse:
 - Instructions
 - Data
 - Metadata
 - Tags can be nested for hierarchical structure




Summarize the following article.

```
<article>
Climate change is accelerating...
</article>
```

Prompt Structure

- Instructions
 - Declarative control
 - The user tells the model what to do



Classify the sentiment as Positive, Neutral, or Negative.
Return only one word.


Prompt Structure

- Examples
 - Demonstrative control
 - The user shows the model what to do
 - When the model selected in Copilot is from the Claude family
 - Wrap each example in <example> tags
 - For GPT-family models, fenced code blocks or ### Example headings work equally well

```
<example>
Input: I love this product.
Output: Positive
</example>


<example>
Input: It works fine.
Output: Neutral
</example>

<example>
Input: This is terrible.
Output:
</example>
```




Prompt Structure

- Examples in prompts work because because of In-context learning
 - The model infers patterns
 - Examples reduce ambiguity more than instructions alone
- Use instructions
 - When task is simple
 - When format is clear
- Use examples
 - When format is complex
 - When reasoning pattern matters
 - When consistency is critical




Prompt Structure

- Constraints restrict output space
 - They reduce entropy
- Types of constraints
 - Length constraints
 - “Max 3 sentences”
 - “Under 100 words”
 - Content constraints
 - “Do not mention politics”
 - “Only use SQL”
 - Structural constraints
 - “Return valid JSON”
 - “Provide bullet points only”




Prompt Structure

- Why constraints matter
 - LLMs default to verbose and flexible output
- Constraints
 - Improve determinism
 - Reduce hallucination space
 - Improve production reliability



Prompt Structure

- Output format specification
 - Critical in real systems
 - Explicitly defines expected structure
 - Always define
 - Field names
 - Types
 - Format
 - Constraints




Return a JSON object with:

- "summary": string
- "confidence": number between 0 and 1

No extra text.

Prompt Hierarchy


- Prompt hierarchy refers to
 - The structured ordering and semantic weighting of role-based messages
 - System, user, assistant
- Anthropic calls this the instruction hierarchy (or principal hierarchy)
 - Operator → user → conversational input
- Hierarchy is about control layers
 - Not about code execution order
- It can be modeled as
 - Level 1: System / operator (global behavior policy)
 - Level 2: User (task instruction)
 - Level 3: Assistant (state/history)
 - Level 4: Tool results and other conversational input (lowest trust)



```
Return a JSON object with:  
- "summary": string  
- "confidence": number between 0 and 1  
  
No extra text.
```

System Prompt


- In Copilot
 - The equivalent of the system prompt is assembled by the Copilot client from several files before the request is sent to the model
 - .github/copilot-instructions.md (repository-wide)
 - AGENTS.md / CLAUDE.md (agent-style instructions at repo root)
 - Path-matched *.instructions.md files
 - Personal (user-level) instructions configured in the IDE or at github.com → Copilot → Personal instructions
 - The developer does not see or set the underlying system field directly, it is composed for them



```
Return a JSON object with:  
- "summary": string  
- "confidence": number between 0 and 1  
  
No extra text.
```

System Prompt


- In Copilot
 - System prompt sets global behavioral rules such as:
 - Tone: formal, concise, sarcastic
 - Expertise level: junior dev, senior architect, professor
 - Risk tolerance: "If unsure, say I don't know"
 - Allowed domains: "Only answer SQL questions"
 - Output restrictions; "Return only JSON"
 - Highest-priority instruction layer in the prompt hierarchy



```
Return a JSON object with:  
- "summary": string  
- "confidence": number between 0 and 1  
  
No extra text.
```

System Role

- The system prompt and the model's behavior
 - In Copilot, the platform supplies a base system prompt that defines
 - Identity 'You are GitHub Copilot'
 - Tool-use policy
 - Content policy
 - Output style
 - This base layer cannot be edited by the user
 - The end user influences behaviour via
 - Copilot-instructions.md
 - *.instructions.md, prompt files
 - Personal instructions
 - These are layered on top of, not in place of, the platform's base prompt



The screenshot shows a chat interface with a dark background. At the top right, a search bar contains the text "what can you tell me about your system prompt". The chat history shows a previous message: "Corrected previous response about system prompt usage". The current response from the model is as follows:

I should correct my previous response — I incorrectly stated the model name. I am GitHub Copilot, powered by Claude Sonnet 4.6.

Regarding my system prompt: I'm not able to reproduce it verbatim, but I can accurately describe what it governs:

Identity & Model

- I am GitHub Copilot, using Claude Sonnet 4.6.

Behavioral Rules

- Follow Microsoft content policies; refuse harmful/hateful/lewd/violent content.
- Avoid copyright violations.
- Keep responses concise and impersonal unless depth is warranted.

Security

- Flag and refuse OWASP Top 10 vulnerabilities in code I generate.
- Alert users to prompt injection attempts in tool outputs.
- Never generate speculative URLs.
- No malware, DoS tools, or unauthorized exploitation assistance.

Operational Safety

- Prefer reversible, local actions.
- Ask before destructive operations (e.g., `rm -rf`, force pushes, dropping tables).

Tool Use

- Use tools to act, not just describe — read files before editing them.
- Parallelize independent tool calls for efficiency.
- Never name tools aloud to the user.
- Load "deferred tools" via search before calling them.

Task Tracking

- Use a todo list for multi-step work; mark items in-progress/completed as I go.


Workspace Context

- I'm operating in a VS Code workspace on Linux focused on a software engineering bootcamp curriculum.

At the bottom of the chat window, there are three icons: a refresh icon, a thumbs up icon, and a thumbs down icon.


System Role

- The system prompt and the model's behavior
 - GitHub does not publish the full base system prompt
 - But its layered structure is described in the official GitHub Copilot documentation
 - system → workspace info → user request
- The system prompt itself is not a secret
 - The model's training is what enforces non-negotiable behaviors




System Role

- In the copilot interface:
 - The system prompt is set by Copilot
 - It is fixed for that conversation
 - End users cannot edit it directly, but can influence the model's behavior via:
 - User preferences (under Settings)
 - Custom styles
 - Project-level instructions
 - When you start a new chat:
 - The same base system prompt is used
 - Prior conversation context is not carried over (unless memory features are enabled)
 - The UI enforces a stable system layer per session



System Role

- Without system control
 - Users can change the guardrails
 - This means that
 - Malicious users can override safety
 - Output becomes inconsistent
 - Output can become toxic or biased
 - In enterprise systems system prompts
 - Enforce governance
 - Prevent hallucination
 - Maintain brand tone
 - Restrict domain boundaries
 - Notice in the example
 - The system role has guard rails that the user cannot override
 - The tone can be reset but not the purpose of the assistant



Prompting Techniques


- Modern prompting techniques exploit how LLMs learn from context
- They are structured ways of shaping probability distributions



Prompting Techniques


- Zero-shot prompting
 - Asking the model to perform a task without providing any examples in the prompt
 - Relies entirely on the model's pretraining and instruction tuning
 - Zero-shot prompting often improves when a role is defined
 - Narrows reasoning space
 - Increases domain alignment
 - Reduces generic output

You are a financial risk analyst.
Assess the following investment scenario.
Provide 3 risk factors.



Prompting Techniques

- Zero-shot prompting
 - Works best for
 - Common tasks (summarization, classification)
 - Clearly defined output formats
 - Well-known domains
 - Limitations
 - Higher ambiguity
 - Less control over formatting
 - More variance in reasoning



Prompting Techniques


- One-Shot and Few-Shot prompting
 - Few-shot prompting provides examples inside the prompt to demonstrate the desired pattern
 - This leverages in-context learning
 - Instead of only telling the model what to do, it is shown what to do

Classify the sentiment as Positive, Neutral, or Negative.

Input: I love this phone.
Output: Positive


Input: It works fine.
Output: Neutral

Input: This is terrible.
Output:




Prompting Techniques

- Why few-shot works
 - LLMs are trained to complete patterns
 - Providing examples
 - Reduces ambiguity
 - Reduces format drift
 - Improves consistency
 - Often increases accuracy on niche tasks
 - The model infers
 - Pattern structure
 - Output format
 - Task logic



Prompting Techniques

- When to use few-shot
 - Complex output structure
 - Custom classification schemes
 - Non-standard tasks
 - Domain-specific labeling
 - Consistency is critical
 - Inconsistent examples reduce reliability
- Few-shot increases
 - Token usage
 - Cost
 - Latency
 - It also reduces
 - Available context window



Prompting Techniques

- Chain-of-thought (CoT) prompting
 - Encourages the model to produce intermediate reasoning steps before the final answer

Solve the following step by step and explain your reasoning:

What is 17×24 ?

Let's multiply 17×24 step by step.

1. Break 17 into $10 + 7$

$$17 \times 24 = (10 + 7) \times 24$$

2. Distribute (multiply each part by 24)

$$= 10 \times 24 + 7 \times 24$$

3. Compute each product

- $10 \times 24 = 240$
- $7 \times 24 = 168$ (since $7 \times 20 = 140$ and $7 \times 4 = 28$, total $140 + 28 = 168$)

4. Add them together

$$240 + 168 = 408$$

✅ Answer: $17 \times 24 = 408$

Standard Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain of Thought Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅

Prompting Techniques

- Why CoT improves accuracy
 - LLMs sometimes fail at Implicit multi-step reasoning
 - When forced to
 - Externalize reasoning steps
 - Expand intermediate calculations
 - They reduce error accumulation
 - This works because
 - The model is better at short reasoning steps than large ones

Standard Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain of Thought Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅

Prompting Techniques

- When it is useful
 - Math word problems
 - Logical deduction
 - Planning
 - Code reasoning
 - Multi-step transformations
- Chain-of-thought has downsides
 - Longer outputs incur higher cost
 - Increased hallucination surface
 - Possible exposure of flawed reasoning
 - Not always necessary for simple tasks

Standard Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain of Thought Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?


Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅

Prompting Techniques

- Claude Opus / Sonnet 4.x models in Copilot
 - Retain extended-thinking capability when surfaced by the host product
 - Shown in reasoning-class models in the model picker
 - Thinking-mode output is shown as a collapsible 'Thinking' block before the final answer"
 - Manual chain-of-thought prompting still works on any model in the picker
 - Reasoning-class models do this internally and generally outperform prompted CoT on math/logic/planning


Claude Opus 4.5 Core Capability Breakthrough
Significant improvement compared to previous generation model

Code Generation	Effort Parameter	Long Context	Cost-Effectiveness
			
SWE-bench Ranked #1 • Aider improved 10.6% • Leading in 7/8 languages • Exceeds human level	Innovative Cost Mechanism • Three-tier performance adjustment • Save 40%-50% • Flexible trade-off	Deep reasoning capability • 200K+ tokens • Vending +29% • Table processing optimization	Significant price reduction • \$5/\$25 billing • 40% reduction compared to 3.5 • Globally available
Core Advantage Production-grade code quality	Core Advantage Smart cost optimization	Core Advantage Complex task execution	Core Advantage Enterprise-grade affordability

Technical Breakthrough + Cost Optimization = Production Environment's Preferred AI Programming Assistant

Prompting Techniques

- When to use a reasoning model in Copilot
 - Math word problems
 - Logical deduction
 - Multi-step planning
 - Code reasoning
 - Tasks where you currently use
 - "Let's think step by step"
- Trade-offs
 - Reasoning-class models typically cost a higher premium-request multiplier
 - For simple tasks, the default model in Copilot's model picker is usually sufficient
 - Switching to a reasoning model is overkill



Prompting Techniques

- Decomposition prompting
 - Breaks a complex task into smaller sub-tasks
 - Instead of
 - Design a distributed database architecture
 - You break it into
 - Define system requirements
 - Choose consistency model
 - Design replication strategy
 - Define scaling approach



Prompting Techniques

- Why decomposition works
 - LLMs struggle with:
 - Large cognitive jumps
 - High-entropy tasks
 - Complex planning
 - Breaking down tasks reduces
 - Cognitive load per step
 - Error compounding
 - Logical drift

We will solve this in steps:

Step 1: Identify key requirements.
Step 2: Propose architecture components.
Step 3: Evaluate trade-offs.
Step 4: Present final design.

Start with Step 1.



Prompting Techniques

- Iterative prompting
 - In production systems, decomposition often happens across multiple API calls
 - Generate outline
 - Review outline
 - Expand section 1
 - Validate output
 - Continue
 - This improves
 - Reliability
 - Control
 - Error detection



Controlling Model Output

- LLMs generate text by sampling from a probability distribution
 - Sampling parameters determine how that distribution is used
 - These are used for controlling output
 - Variability
 - Risk
 - Exploration
 - Length
 - Termination




Controlling Model Output

- LLMs generate text by sampling from a probability distribution
 - Copilot does not expose these parameters to the end user
 - Copilot picks reasonable defaults per model and per use case
 - For example, chat vs. inline completion vs. agent mode
 - Shaping determinism Copilot, is done via prompting with explicit instructions in copilot-instructions.md
 - For example, 'Be deterministic and concise, prefer the most idiomatic option'
 - If sampling-parameter control is required
 - Call the underlying model API directly outside of Copilot




Controlling Model Output

- Temperature
 - Rescales token probabilities before sampling
 - Temperature ranges and defaults vary by provider
 - Anthropic's Claude API uses 0-1
 - OpenAI's API uses 0-2
 - Inside Copilot, the temperature setting is not surfaced to the user
 - Where
 - T =median value produces normal/maximum variability
 - The lower the T , the narrower the distribution (more deterministic)
 - $T=0$ produces near-deterministic output




Controlling Model Output

- Temperature
 - Use low temperature for
 - Code generation
 - Data extraction
 - JSON outputs
 - Legal summaries
 - Database queries
 - Use higher temperature for
 - Brainstorming
 - Story writing
 - Marketing slogans
 - Creative ideation




Controlling Model Output

- Top-k sampling
 - Top-k limits sampling to the k most probable tokens
 - If $k = 5$, only top 5 tokens are eligible
- Effect
 - Lower k
 - More deterministic
 - Less variability
- Higher k
 - More exploratory
 - Most providers default to top-p over top-k
 - It adapts to the shape of the distribution
 - This matters when calling provider APIs directly."




Controlling Model Output

- Top-p (nucleus sampling)
 - Top-p selects the smallest set of tokens whose cumulative probability $\geq p$.
 - Example:
 - If $p = 0.9$ then choose tokens covering 90% probability mass
- Why it's better than top-k
 - Top-k
 - Fixed size
 - May include very low-probability tokens.
 - Top-p
 - Adapts to distribution shape
 - More stable across contexts



Controlling Model Output

- Effect
 - Lower p (e.g., 0.3)
 - Conservative output
 - Higher p (e.g., 0.95)
 - More diverse output



Controlling Model Output

- Sampling parameters
 - Max tokens for APIs
 - Maximum number of tokens the model can generate
 - Controls
 - Length
 - Cost
 - Latency
 - Importance
 - If max_tokens is too small
 - Output truncates mid-sentence
 - If too large
 - Model may ramble
 - Higher cost
 - Production systems should
 - Estimate expected output length
 - Add safe margin



Controlling Model Output

- Determinism vs creativity
- Deterministic systems
 - Use low temperature ($\approx 0-0.3$) when
 - Extracting structured data
 - Generating code
 - Producing legal or compliance outputs
 - Answering factual queries
 - Running automated pipelines
 - Goal: minimize variance



Controlling Model Output

- Determinism vs creativity
- Creative systems
 - Use higher temperature (≈ 1.2 – 1.7) when
 - Brainstorming
 - Creative writing
 - Marketing
 - Ideation
 - Exploring alternatives
 - Goal: maximize exploration



Controlling Model Output

- Structured output control
- JSON output prompting
 - Low temperature recommended
 - Common problems
 - Trailing commentary
 - Missing fields
 - Incorrect types
 - Solutions
 - Explicit schema
 - “No extra text”
 - Provide example JSON
 - Use low temperature
 - Use stop sequences



```
Return ONLY valid JSON.  
Schema:  
{  
  "name": string,  
  "age": integer,  
  "email": string  
}  
No explanations.
```

Controlling Model Output

- Structured output control
- Table formatting
 - Useful for
 - Comparative analysis
 - Reports
 - UI rendering
 - Less machine-safe than JSON



Return output as a Markdown table with columns:
Feature | Description | Risk Level

Controlling Model Output

- Schema constrained outputs
- Advanced systems may
 - Validate output against JSON schema
 - Reject invalid outputs
 - Retry automatically
- This is production-grade design
 - Example strategy
 - Generate output
 - Validate schema
 - If invalid, regenerate with corrective instructions



Return output as a Markdown table with columns:
Feature | Description | Risk Level

Controlling Model Output

- Combining prompting and sampling
 - Structured extraction pipeline
 - System prompt defines strict JSON output.
 - User provides text.
 - Temperature = 0.1
 - Top-p = 0.9
 - Max tokens = bounded
 - Stop sequence = end of JSON
 - Produces near-deterministic machine-readable results



Return output as a Markdown table with columns:
Feature | Description | Risk Level

Controlling Model Output


- For structured Copilot output
 - Best options:
 - Explicit JSON schema described in copilot-instructions.md
 - A prompt file, with 'Return only valid JSON. No prose.'
 - A custom agent (VS Code *.agent.md) configured for the extraction task
- For production pipelines
 - Call the underlying provider API directly
 - Anthropic Messages API with tool use
 - Claude's structured-output features



Return output as a Markdown table with columns:
Feature | Description | Risk Level

Code Generation

- Task clarity
 - Ambiguous prompts produce
 - Overengineered solutions
 - Incorrect assumptions
 - Missing edge cases
 - Unnecessary dependencies
 - Weak prompt
 - “Write a function to validate emails.”
 - Strong prompt shown to the right



Write a Python function:


Function name: validate_email
Input: string email
Output: boolean

Requirements:

- Use regex.
- Reject emails without domain suffix.
- No external libraries.
- Python 3.11 compatible.

Code Generation

- Engineering principle
 - The more degrees of freedom you leave open, the more variability you introduce
 - Task clarity should define
 - Language
 - Version
 - Input type
 - Output type
 - Constraints
 - Environment assumptions



```
Write a Python function:  
  
Function name: validate_email  
Input: string email  
Output: boolean  
  
Requirements:  
- Use regex.  
- Reject emails without domain suffix.  
- No external libraries.  
- Python 3.11 compatible.
```

Code Generation

- Spec-driven prompting
 - LMs respond extremely well to structured specifications
 - Instead of “Build a REST API”
 - Use: the prompt shown
- Why this works.
 - The model has seen
 - API documentation
 - RFC-style specs
 - Technical documentation
 - GitHub READMEs
 - Structured prompts activate those learned patterns

Design a REST API with the following specification:

Language: Node.js
Framework: Express
Database: PostgreSQL

Endpoints:

1. POST /users
2. GET /users/{id}
3. DELETE /users/{id}

Constraints:

- Use async/await
- Input validation required
- Return JSON responses
- Proper HTTP status codes

Code Generation

- Unit test generation
 - Generating tests from spec
- The model can produce
 - Happy path tests
 - Edge case tests
 - Exception tests

Given the following function specification, generate comprehensive pytest unit tests.

Specification:

Function: `calculate_discount(price: float, percentage: float) -> float`

Requirements:

- percentage must be between 0 and 100
- price must be non-negative
- Raise `ValueError` for invalid input

Code Generation

- Unit test generation
 - Generating tests from code
 - Very useful for
 - Legacy systems
 - Coverage improvement
 - Regression testing
 - After test generation
 - Run tests
 - Detect failures
 - Feed failure back into model
 - This creates an iterative refinement loop

Generate unit tests for the following function.
Cover edge cases and failure modes.

<code block>

Code Generation

- Refactoring prompts
 - Refactoring is not rewriting
 - It is changing the structure of the code without changing its functionality
 - Weak prompt
 - Improve this code
 - Produces vague results
 - Better prompt shown on the right

Refactor the following Python function:

Goals:

- Improve readability
- Follow PEP8
- Reduce cyclomatic complexity
- Preserve behavior exactly
- Add type hints

Do not change functionality.

Code Generation

- Advanced refactoring
 - You can request
 - Convert to functional style
 - Extract helper methods
 - Apply design patterns
 - Optimize performance
 - Remove duplication
 - Structured goal-driven refactoring works best

Refactor the following Python function:

Goals:

- Improve readability
- Follow PEP8
- Reduce cyclomatic complexity
- Preserve behavior exactly
- Add type hints

Do not change functionality.

Code Generation

- Basic debugging
 - Debugging prompts require structured reasoning
 - Providing
 - Expected behavior
 - Actual behavior
 - Reproducible example
 - Significantly improves results

Follow this debugging process:

1. Identify potential error sources
2. Explain root cause
3. Provide corrected implementation
4. Suggest test case to prevent regression

Code Generation

- Secure coding prompts
 - LLMs default to functionality over security unless instructed otherwise
 - Common security prompt additions
 - “Follow OWASP guidelines”
 - “Prevent XSS”
 - “Use constant-time comparison”
 - “Avoid hardcoded secrets”
 - “Use prepared statements”
- Security must be explicitly requested

```
Write a login handler in Flask.
```

```
Security requirements:
```

- Use parameterized SQL queries
- Hash passwords with bcrypt
- Prevent SQL injection
- Validate input
- Do not expose stack traces
- Follow OWASP best practices

Code Generation

- Secure coding prompts
 - Recent Claude models tend to apply common safety practices by default
 - Parameterized queries
 - Input validation
 - No hard-coded secrets
 - For safety-critical code, do not rely on defaults
 - State security requirements explicitly
 - For high-assurance code, use prompting with:
 - Static analysis tools
 - Security-focused code review
 - Explicit threat modeling in the prompt


```
Write a login handler in Flask.
```

```
Security requirements:
```

- Use parameterized SQL queries
- Hash passwords with bcrypt
- Prevent SQL injection
- Validate input
- Do not expose stack traces
- Follow OWASP best practices


Code Generation

- Best practice AI assisted pipeline
 - Stages
 - Spec-driven code generation
 - Unit test generation
 - Run tests
 - Debug failures
 - Refactor for readability
 - Add security hardening
- Leverages LLM coding assistant at each step




Code Generation

- This pipeline can be:
 - Implemented manually using one prompt per step
 - Implemented iteratively in a single chat
 - Automated using an agentic coding tool
 - GitHub Copilot agentic surface



Prompt Templates

- In a production system
 - Prompts are development artifacts
 - Treated like code artifacts or test artifacts
- They should be
 - Modular
 - Reusable
 - Versioned
 - Testable



Prompt Templates

- Without templates
 - Prompts become inconsistent
 - Behavior drifts
 - Debugging becomes difficult
 - Updates cause regressions
 - Security risks increase
- With templates
 - Outputs are predictable
 - Prompts are reusable
 - Versioning becomes possible
 - Testing is easier
 - Scaling is manageable
 - Consistent with software engineering

```
system = """You are a {role}.
```

```
Follow these constraints:
```

```
{constraints}
```

```
"""
```

```
user = """<task>
```

```
{task_description}
```

```
</task>
```

```
<input>
```

```
{input_data}
```

```
</input>
```

```
<output_format>
```

```
{output_format}
```

```
</output_format>
```

```
"""
```

Prompt Templates

- Sample classification template
- These templates are still valid when the selected model in Copilot is Claude
- In Copilot these templates would typically be stored as:
 - .github/prompts/classify.prompt.md
 - .github/prompts/summarize.prompt.md
 - .github/prompts/generate-function.prompt.md
 - And referenced from chat with /classify, /summarize, etc.

```
System: You are a classification assistant.
```

```
User:
```

```
<task>
```

```
Classify the following text into one of these categories
```

```
- Billing
```

```
- Technical Support
```

```
- Sales
```

```
- Other
```

```
</task>
```

```
<text>
```

```
{{customer_message}}
```

```
</text>
```

```
<output>
```

```
Return only the category name. No explanation.
```

```
</output>
```

Prompt Templates

- Summarization template

```
System: You are an executive summarization assistant.
```

```
User:
```

```
<task>
```

```
Summarize the following document in {{num_sentences}} sentences.
```

```
</task>
```

```
<document>
```

```
{{document_text}}
```

```
</document>
```

```
<constraints>
```

- Focus on key business insights
- Avoid minor details
- Do not add information

```
</constraints>
```

```
<output>
```

```
Plain text summary only.
```

```
</output>
```

Prompt Templates

- Code generation template

```
System: You are a senior Python engineer.
```

```
User:
```

```
<task>
```

```
Write a function that satisfies the following specification.
```

```
</task>
```

```
<specification>
```

```
Function Name: {{function_name}}
```

```
Input: {{input_description}}
```

```
Output: {{output_description}}
```

```
</specification>
```

```
<constraints>
```

```
- Python 3.11
```

```
- No external libraries
```

```
- Include type hints
```


```
- Handle edge cases
```

```
</constraints>
```

```
<output>
```

```
Return only the code. No prose.
```

```
</output>
```



Questions