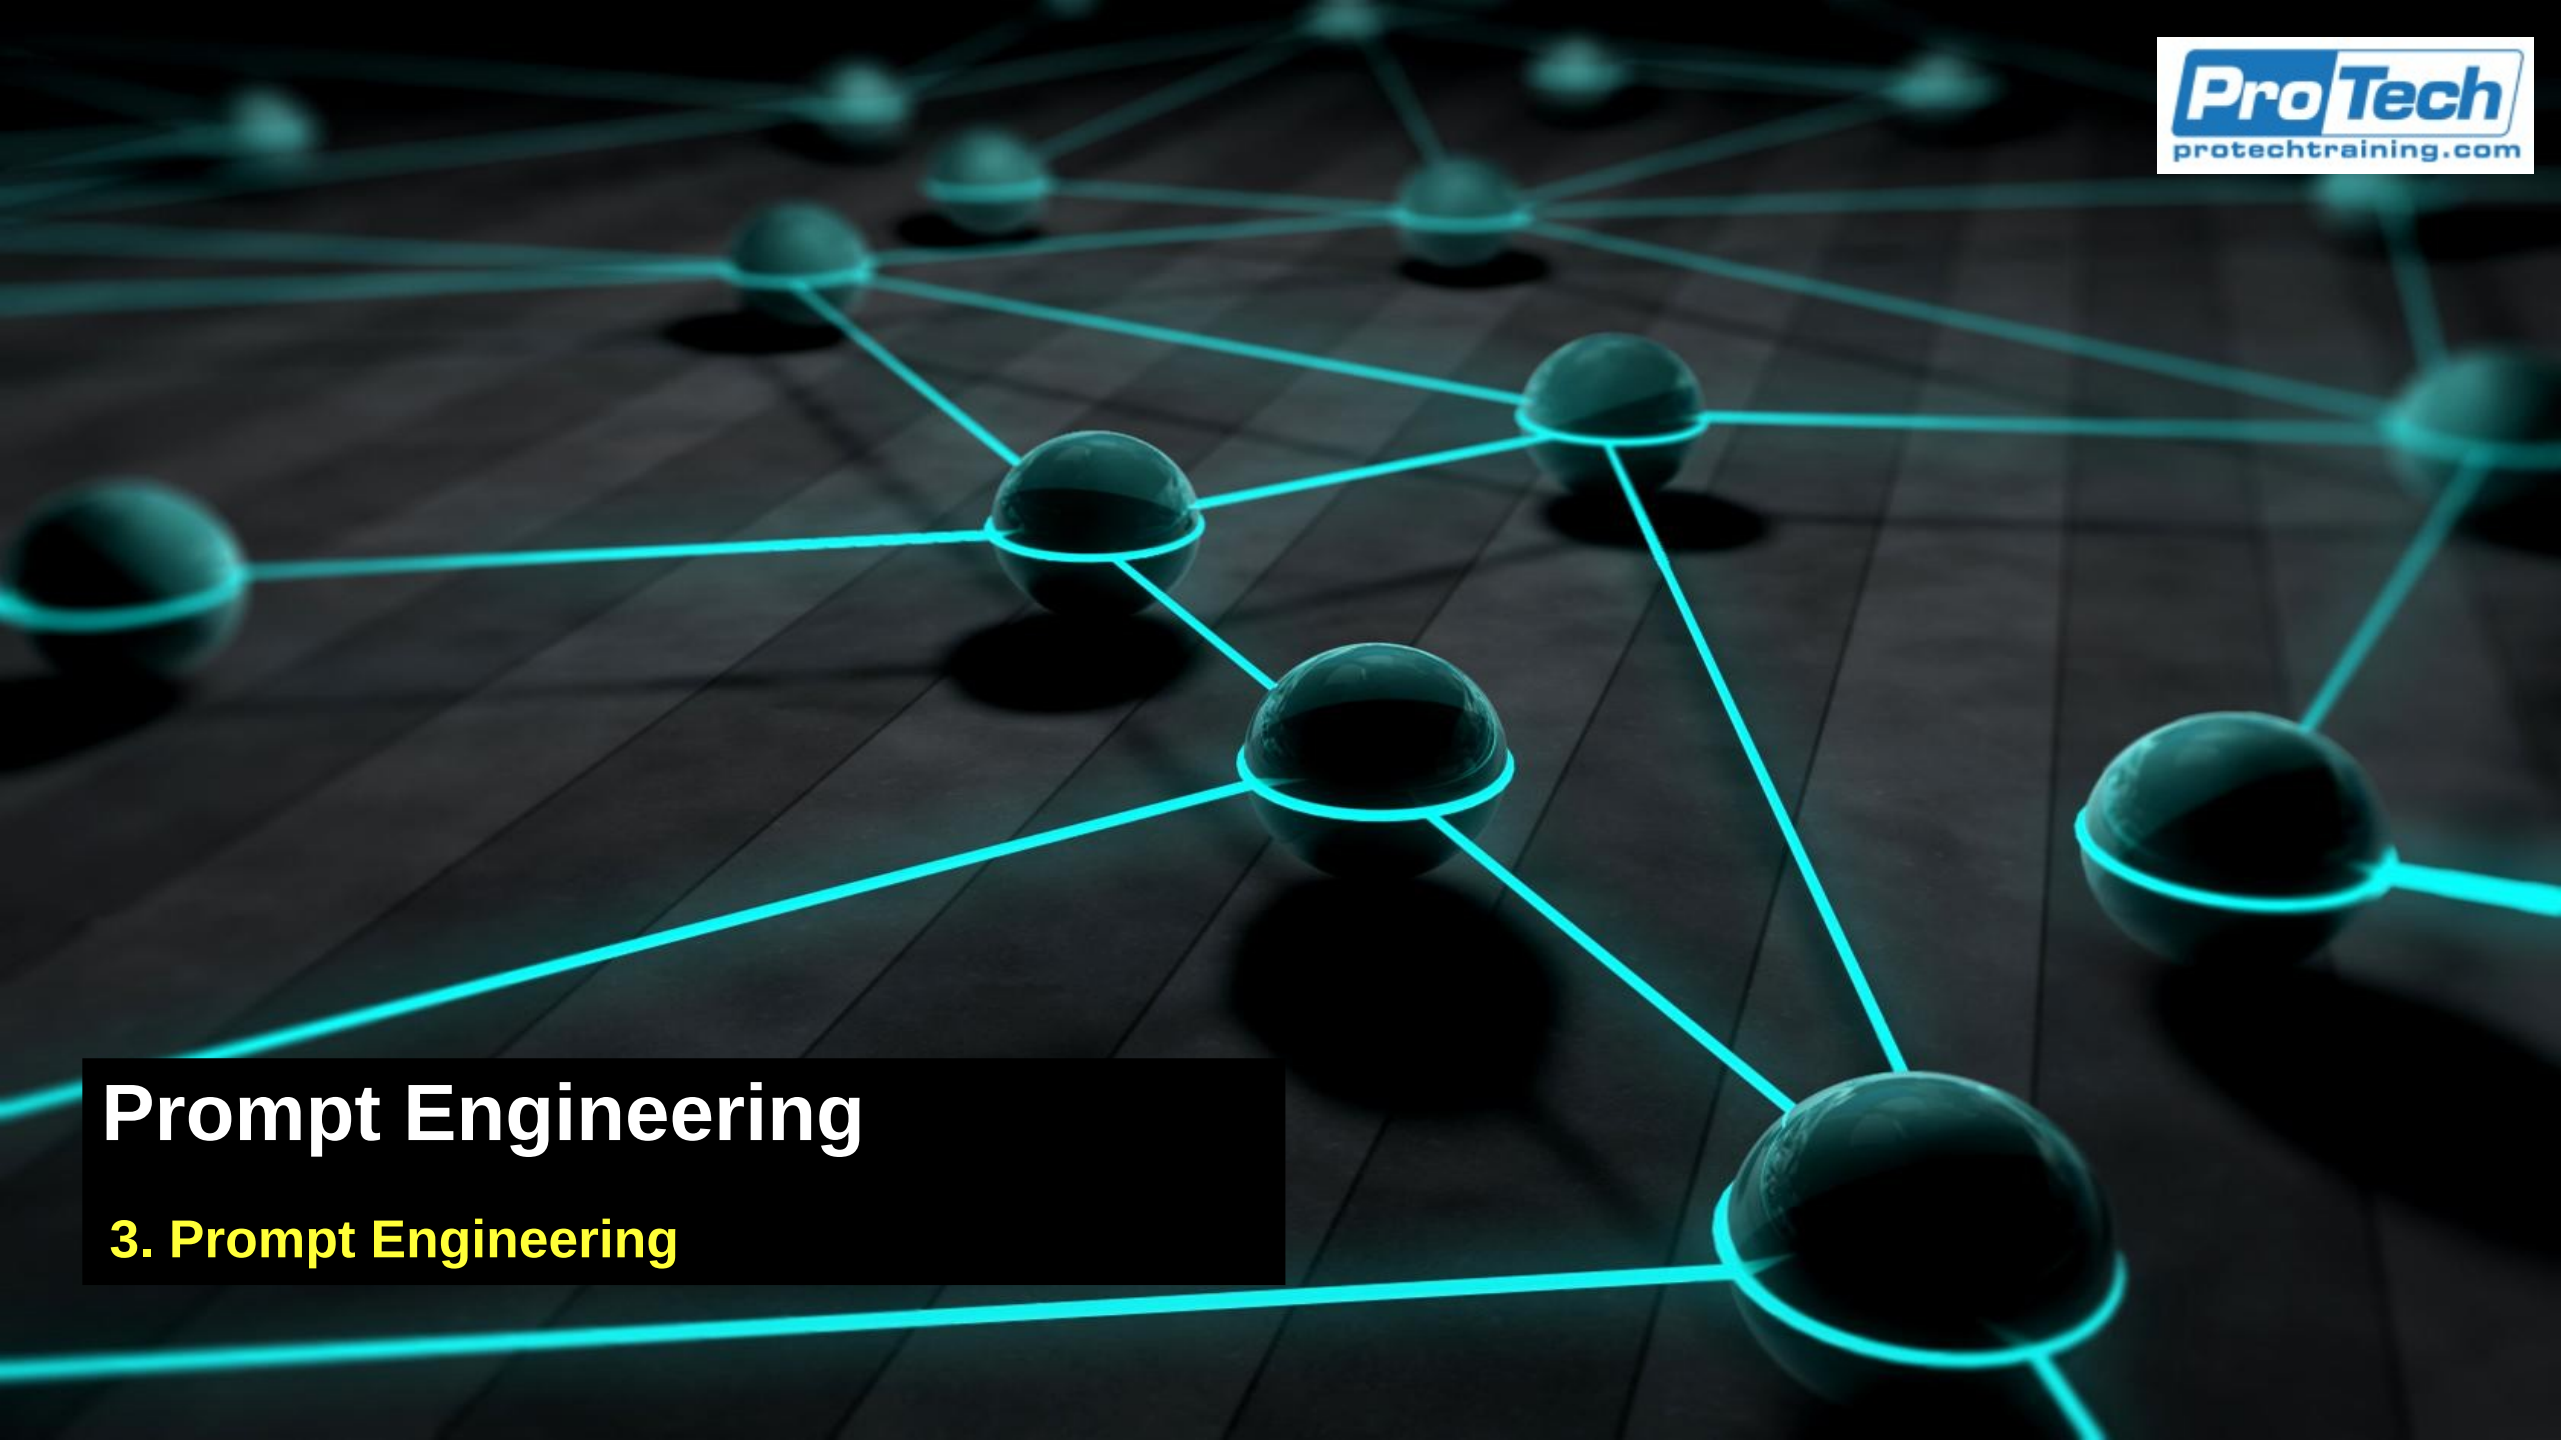

Prompt Engineering

3. Prompt Engineering

Prompt Engineering

- Prompt anatomy
- LLMs don’t “think” like humans, thry
  - Predict the next token
  - Follow patterns
  - Respond strongly to instruction hierarchy
- Well-structured prompts dramatically outperform vague ones

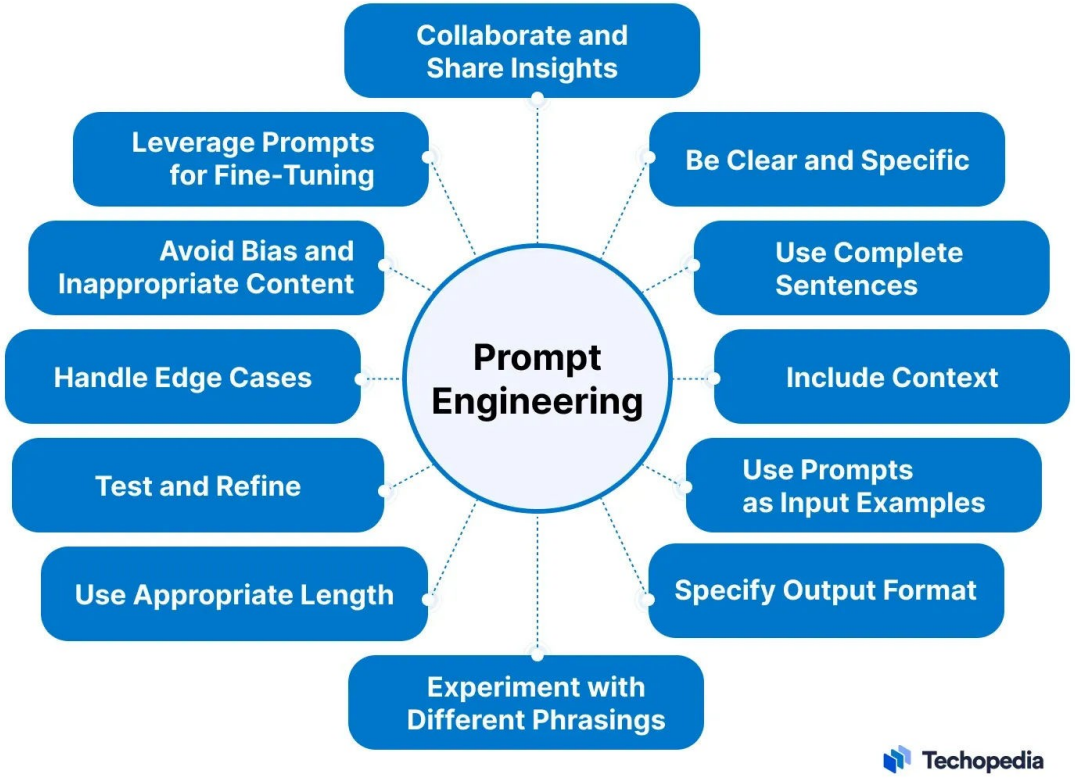

Image Credit: https://botpenguin.com/glossary/prompt-engineering

Prompt Engineering

- Prompt engineering is
  - The systematic design and optimization of input instructions to control and shape the behavior of large language models (LLMs).
  - Prompt engineering = context design for probabilistic model control
- Consists of
  - Structuring instructions
  - Controlling outputs
  - Reducing ambiguity
  - Managing constraints
  - Shaping probability distributions

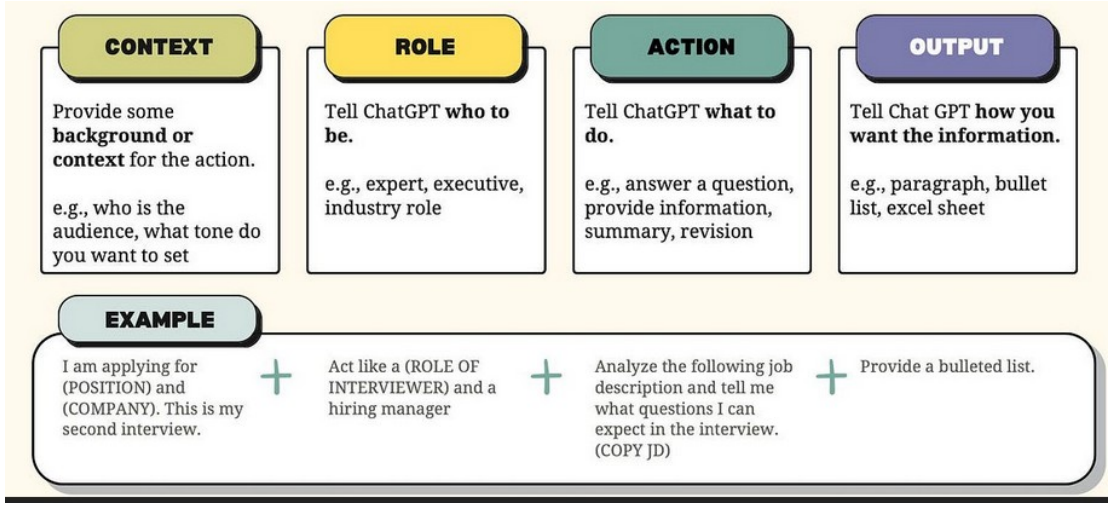

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- LLMs do not execute deterministic programs
  - They compute
    - P(next_token|context)
  - This means that
    - Small changes in input can produce large output shifts
    - Ambiguity leads to unstable outputs
    - Under-specified tasks increase hallucination risk
- Prompt engineering introduces control and predictability

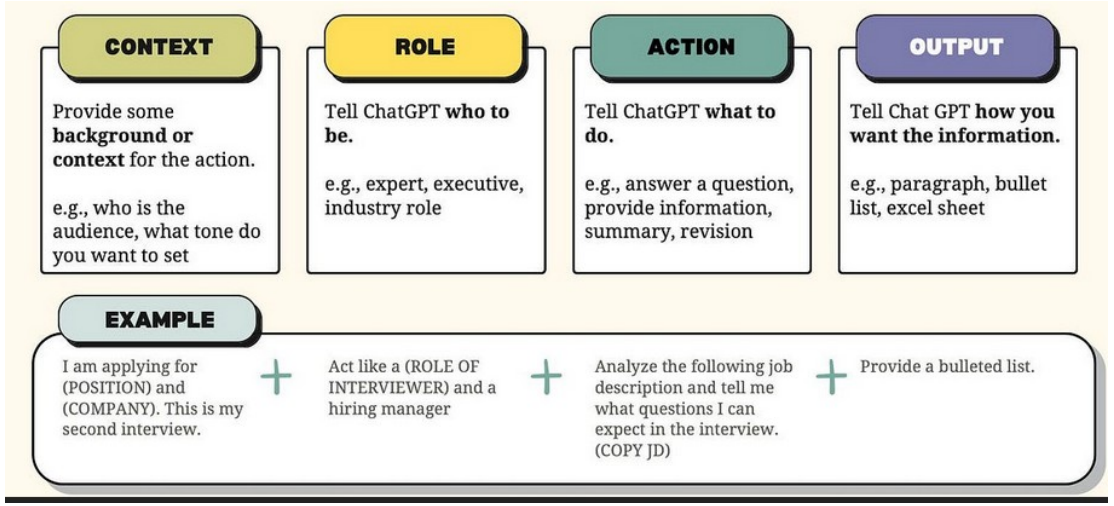

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Prompting works because of how LLMs are trained
- Pretraining
  - Foundation models are trained on massive corpora to learn
    - Language/code structure
    - Patterns
    - Reasoning traces
    - Instruction-response patterns
  - They are not trained for a single task
  - They are trained to approximate
    - P(text_next|text_previous)

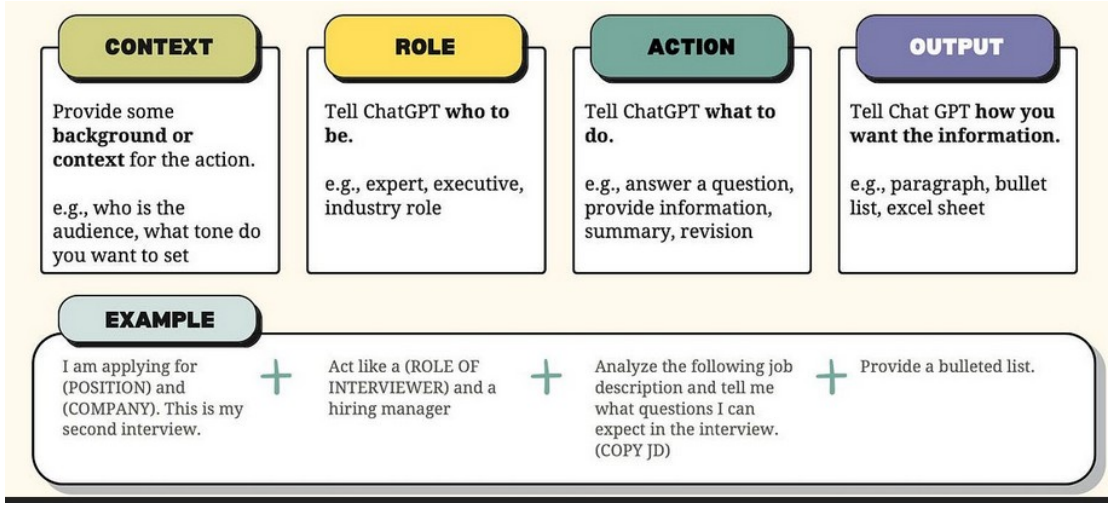

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Instruction tuning
  - Modern models are fine-tuned on
    - Instruction/response pairs
    - Conversations
    - Role-based dialogues
  - This teaches the model that
    - “User message” means requires task fulfillment
    - Structured prompts requires structured output
    - Explicit constraints must be respected
  - The model learns patterns such as
    - If input says “List 3 bullet points”
    - Then output should match that pattern

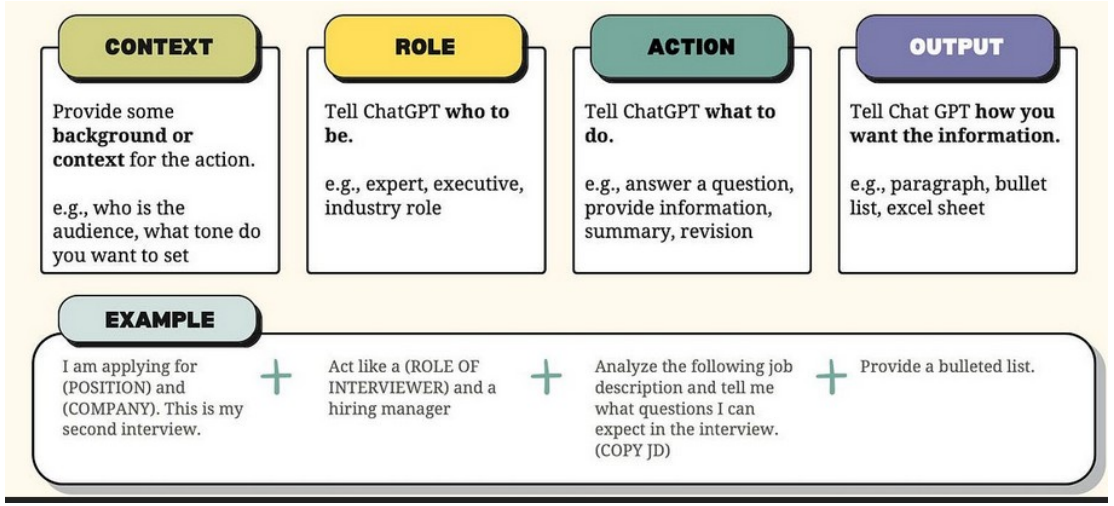

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Prompts as “programming” foundation models
  - Traditional software
    - You write explicit logic
    - The machine executes instructions deterministically
  - With LLMs
    - You describe behavior
    - The model approximates desired outputs probabilistically
  - Prompts act as
  - A high-level control language over a general-purpose stochastic model

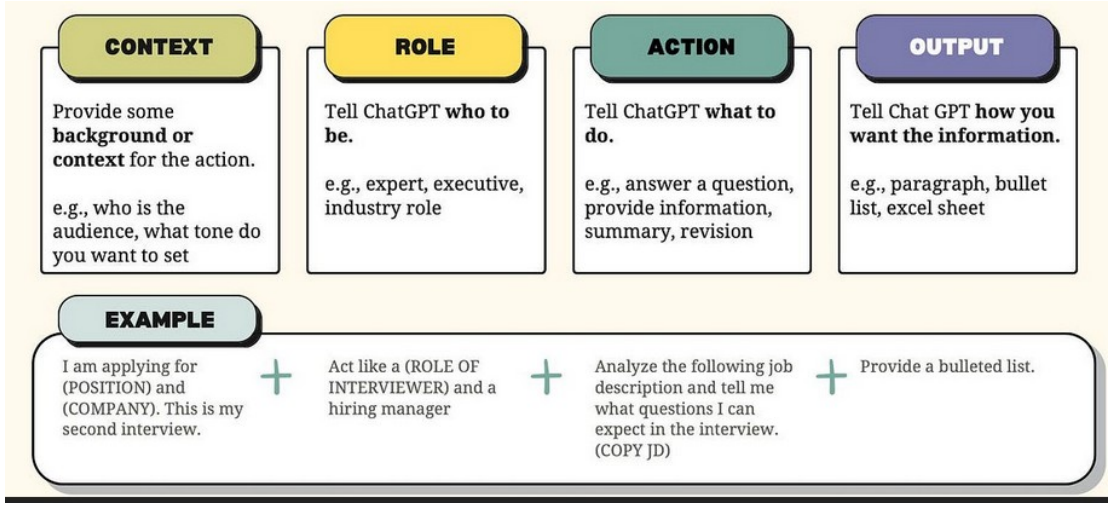

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Programming dimensions in prompts
  - Prompts define:
    - Role
    - Constraints
    - Output format
    - Reasoning style
    - Allowed domain
    - Evaluation criteria
- Example

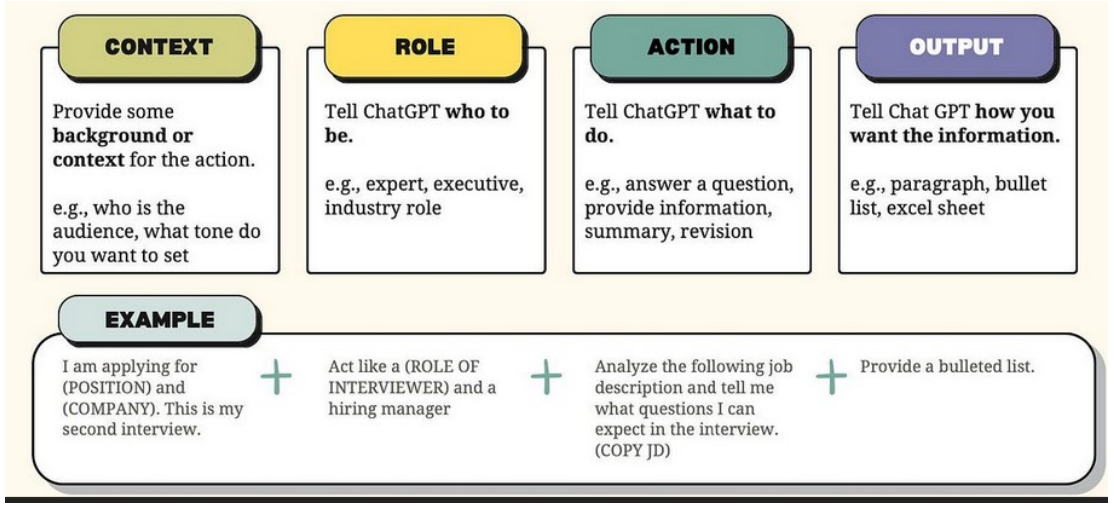


```
You are a strict JSON generator.
Return valid JSON only.
Fields: name (string), age (int).
No extra text.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Limitations of prompting
  - You cannot
    - Install new knowledge permanently
    - Fix core reasoning flaws
    - Remove biases fully
    - Create guaranteed deterministic behavior
  - You are influencing probabilities
    - Not rewriting the model

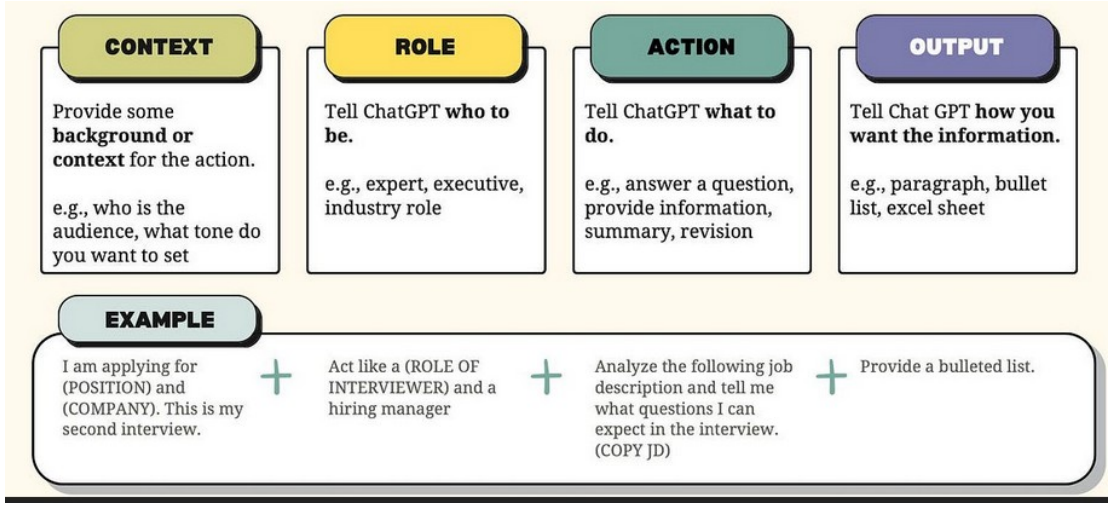

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Limitations of prompting
  - LLMs have finite context (expressed in tokens)
    - That the record of how you have shifted the probability the model uses to generate an answer
    - The actual LLM does not have any form of memory
    - The whole context is uploaded with each request to influence the LLM's respns
- If you exceed the size of the context
  - Earlier instructions are truncated.
  - System prompts may be lost
  - Behavior may drift
- Prompt engineering must consider token budgets

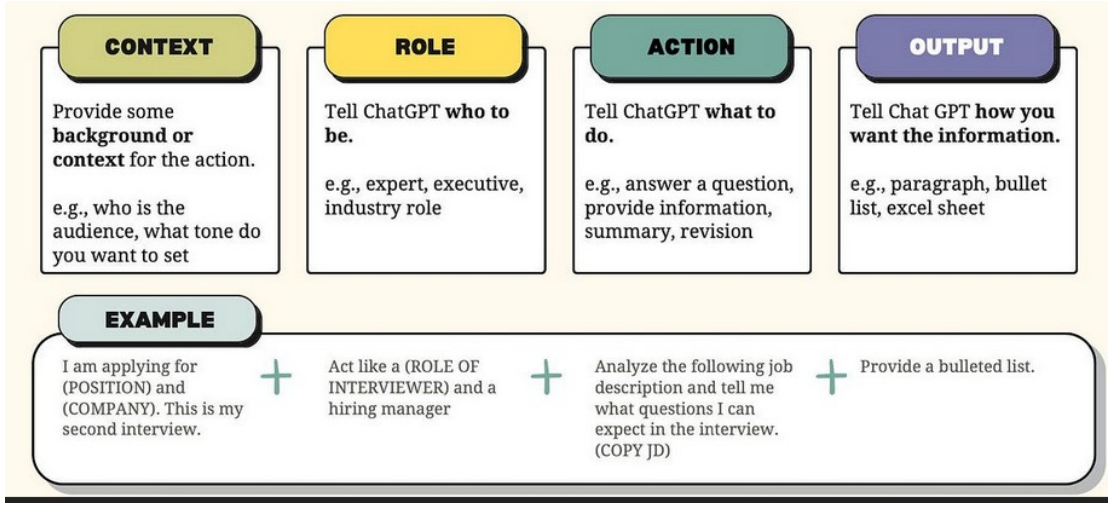

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Limitations of prompting
  - Hallucinations
    - Even with perfect prompting
    - The model may fabricate information
    - It may generate plausible but false outputs
    - Because it predicts likely text, not verified truth
  - Small phrasing differences can cause
    - Output format drift
    - Tone shift
    - Depth variation
    - Constraint violations
- This makes prompting non-deterministic compared to traditional programming

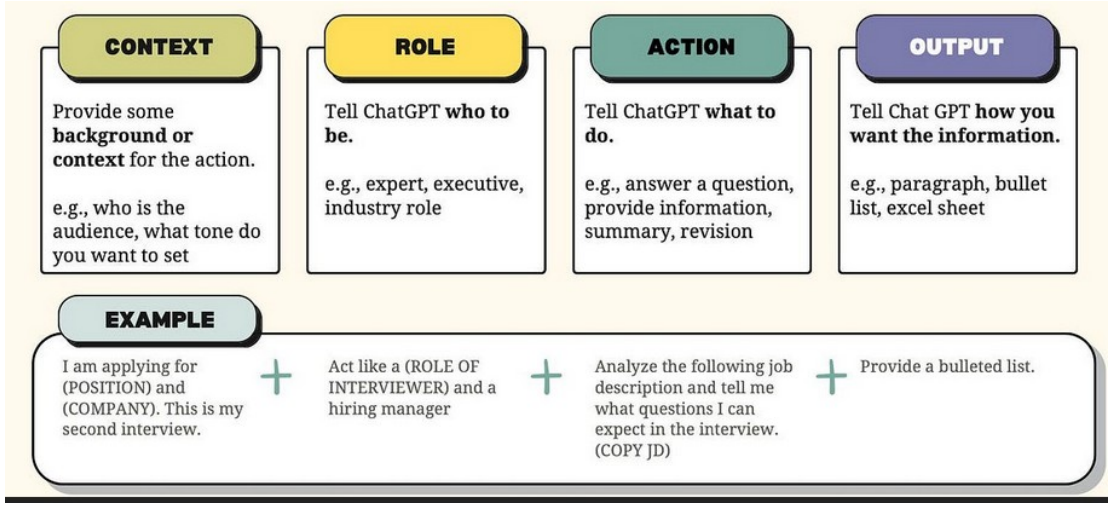

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Engineering

- Complex reasoning limits
  - For problems involving
    - Multi-step symbolic reasoning
    - Formal proofs
    - Long planning chains
  - Prompting alone may not be sufficient
  - You may need
    - Tool use
    - Code execution
    - Retrieval augmentation
    - Structured reasoning frameworks

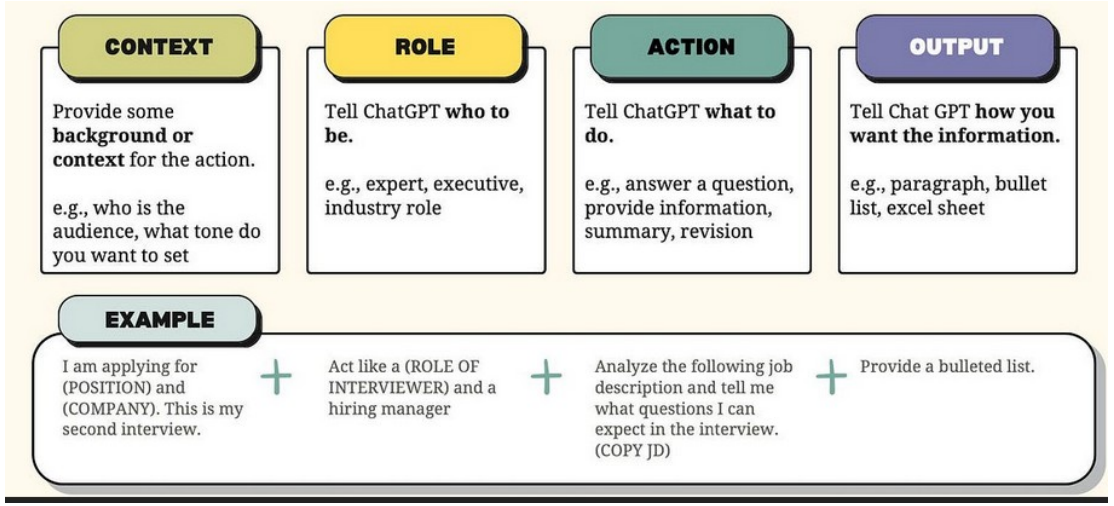

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- A modern LLM call looks like the block on the lower right
  - Inside each message are additional structural components:
    - Context blocks
    - Instructions
    - Examples
    - Constraints
    - Output format specification
  - We can model a full prompt as
    - Role Layer + Task Layer + Constraints +
    - Format Specification + Examples

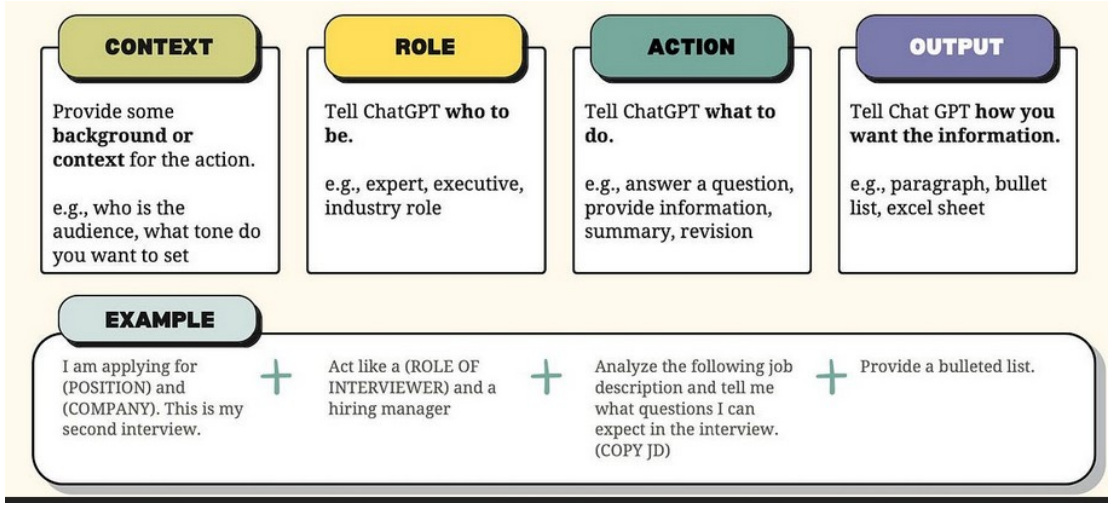


```
Selected model:  Claude Sonnet 4.6   (chosen from VS Code model picker)
Active context:  .github/copilot-instructions.md
                 open editor file + selection
                 workspace metadata (language, framework)
User prompt:     "Refactor this function to use async/await…"
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Copilot Specific
  - Copilot does not expose a raw messages API to the developer
    - The chat surface handles message construction
    - For example: VS Code / Visual Studio / JetBrains / GitHub.com / Copilot CLI
  - The selected model is chosen from the Copilot model picker not via a model field the user sets.
    - Claude Sonnet 4.6, Claude Opus 4.6, GPT-5.x)
  - Behavioral "system prompt" content contributed through
    - github/copilot-instructions.md
    - *.instructions.md
    - And AGENTS.md file
    - Not a system parameter.

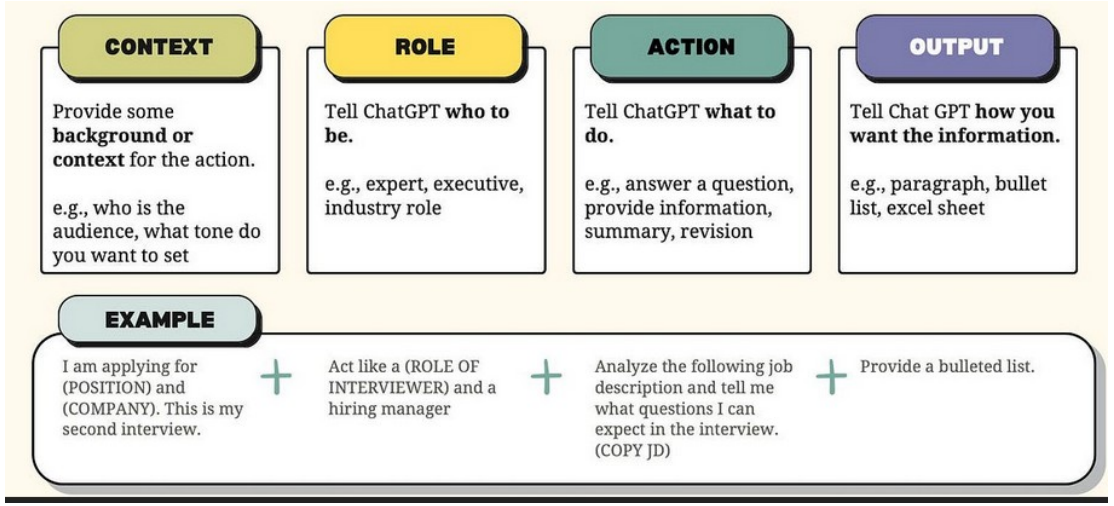


```
Selected model:  Claude Sonnet 4.6   (chosen from VS Code model picker)
Active context:  .github/copilot-instructions.md
                 open editor file + selection
                 workspace metadata (language, framework)
User prompt:     "Refactor this function to use async/await…"
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- System prompt defines
  - Global behavior
  - Identity
  - Policy constraints
- Highest behavioral control layer
  - Contains
    - Persona definition
    - Domain restriction
    - Tone
    - Safety rules
    - Output invariants
    - Formatting policies

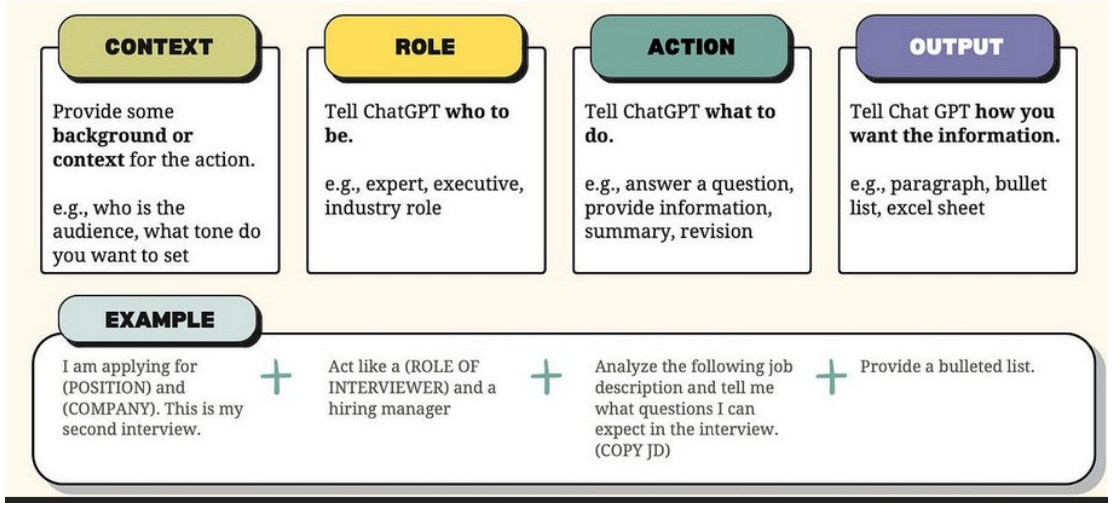


```
You are a senior backend engineer.
Provide concise, technically accurate responses.
If unsure, say "Insufficient information."
Do not invent APIs.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- The user prompt defines
  - The current task or objective
  - It is the instruction-bearing layer
  - The user prompt contains
    - Task request
    - Problem statement
    - Data to analyze
    - Context for current processing step
  - The two prompts can be thought of
    - System prompt defines behavioral rules
    - User prompt defines the task objective

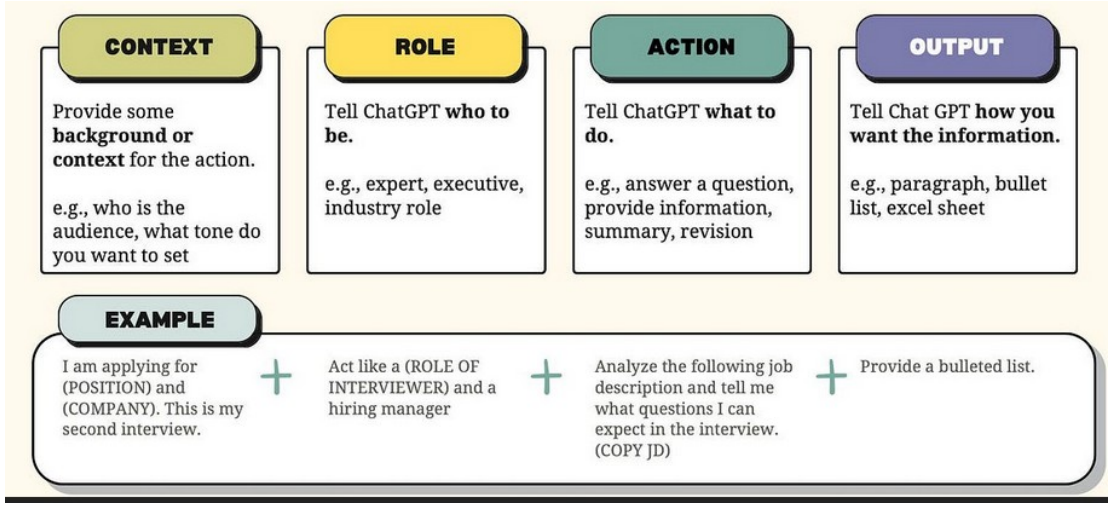


```
Refactor the following Python function to improve readability:

def f(x):
    if x>10: return True
    else: return False
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- The assistant role contains
  - Prior model-generated outputs
  - The contextual state of the conversation
  - Assistant messages
    - Anchor reasoning trajectory
    - Preserve prior decisions
    - Influence consistency
- Example
  - If the assistant previously said
    - *“We’ll solve this using recursion.”*
  - The model is biased to continue recursively

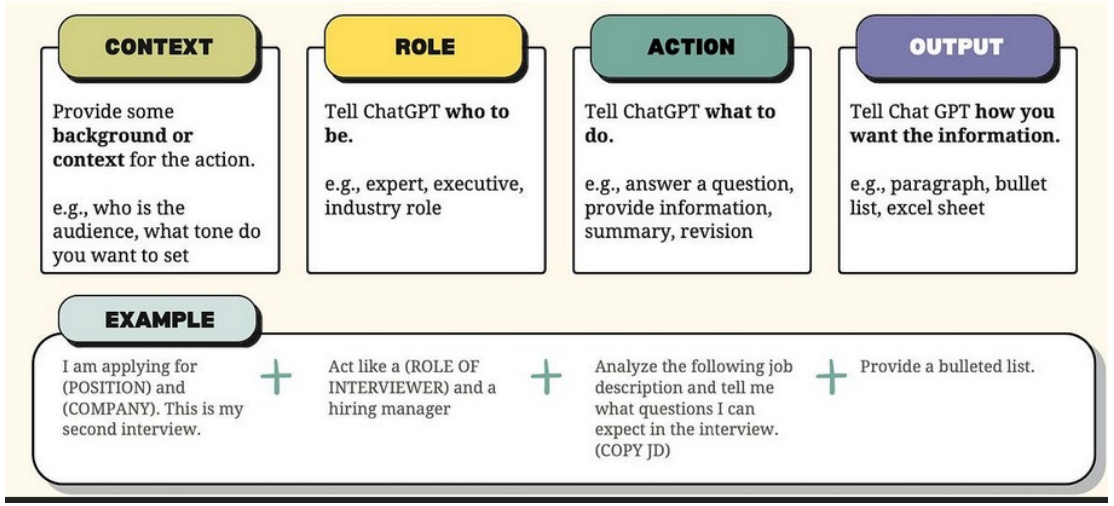


```
Refactor the following Python function to improve readability:

def f(x):
    if x>10: return True
    else: return False
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Context blocks
  - Structured segments inside a message that provide
    - Reference material
    - Data
    - External knowledge
    - Instructions separated from content
  - For Claude, the recommended convention is XML tags as delimiters
    - <article>, <context>, <input>
    - Claude is trained to recognize XML tags and scope content correctly
    - When using a Claude model inside Copilot, the same XML-tag convention applies
  - Without delimiters, the model may confuse:
    - Instructions
    - Data
    - Metadata
  - Tags can be nested for hierarchical structure

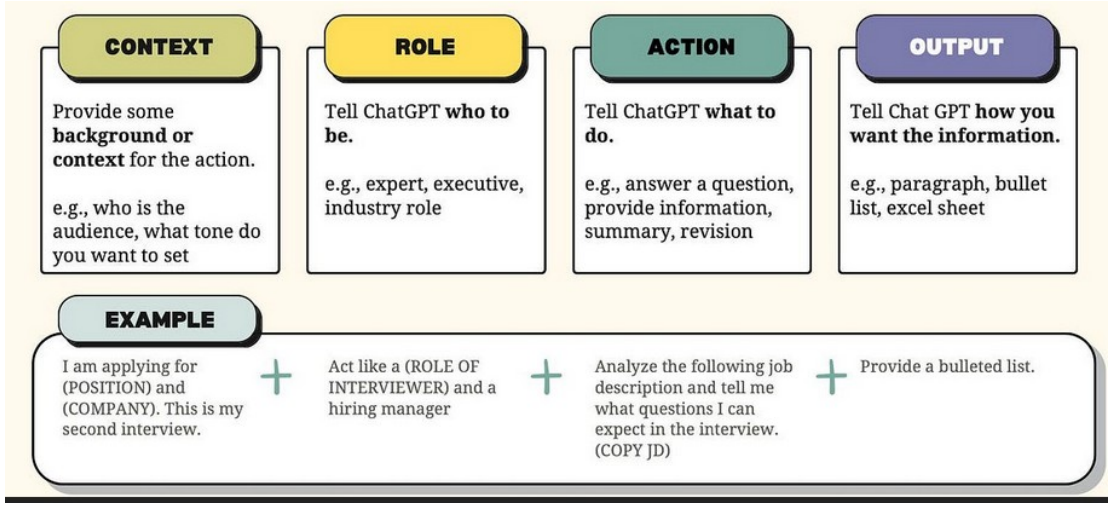


```
Summarize the following article.

<article>
Climate change is accelerating...
</article>
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Instructions
  - Declarative control
  - The user tells the model what to do

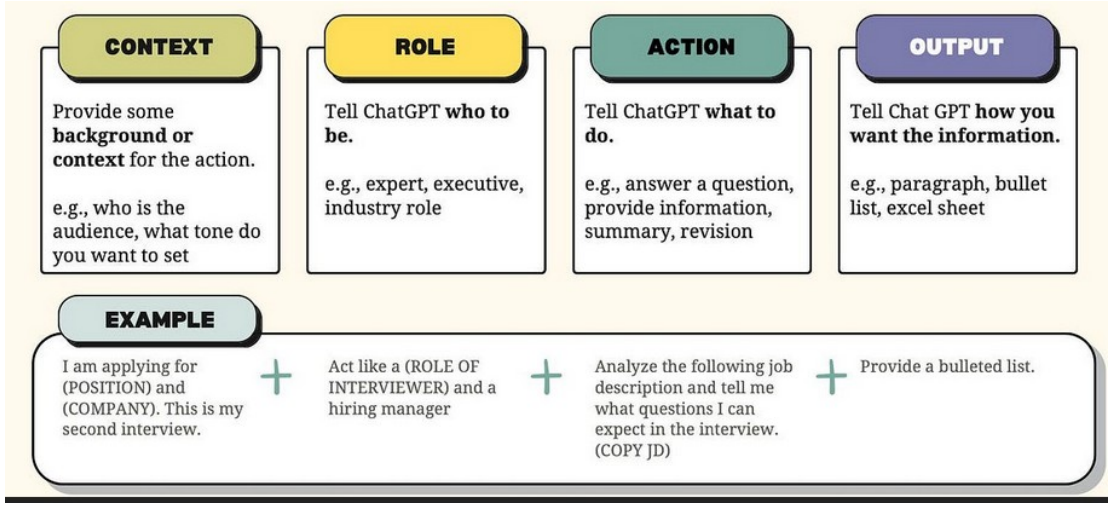


```
Classify the sentiment as Positive, Neutral, or Negative.
Return only one word.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Examples
  - Demonstrative control
    - The user shows the model what to do
  - When the model selected in Copilot is from the Claude family
    - Wrap each example in <example> tags
    - For GPT-family models, fenced code blocks or ### Example headings work equally well

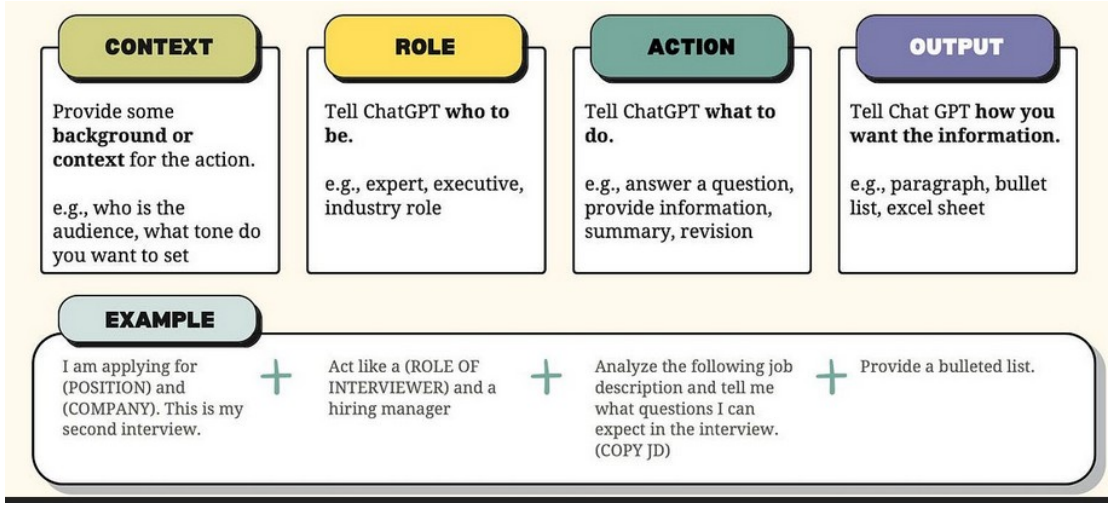


```
<example>
Input: I love this product.
Output: Positive
</example>

<example>
Input: It works fine.
Output: Neutral
</example>

<example>
Input: This is terrible.
Output:
</example>
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Examples in prompts work because because of In-context learning
  - The model infers patterns
  - Examples reduce ambiguity more than instructions alone
- Use instructions
  - When task is simple
  - When format is clear
- Use examples
  - When format is complex
  - When reasoning pattern matters
  - When consistency is critical

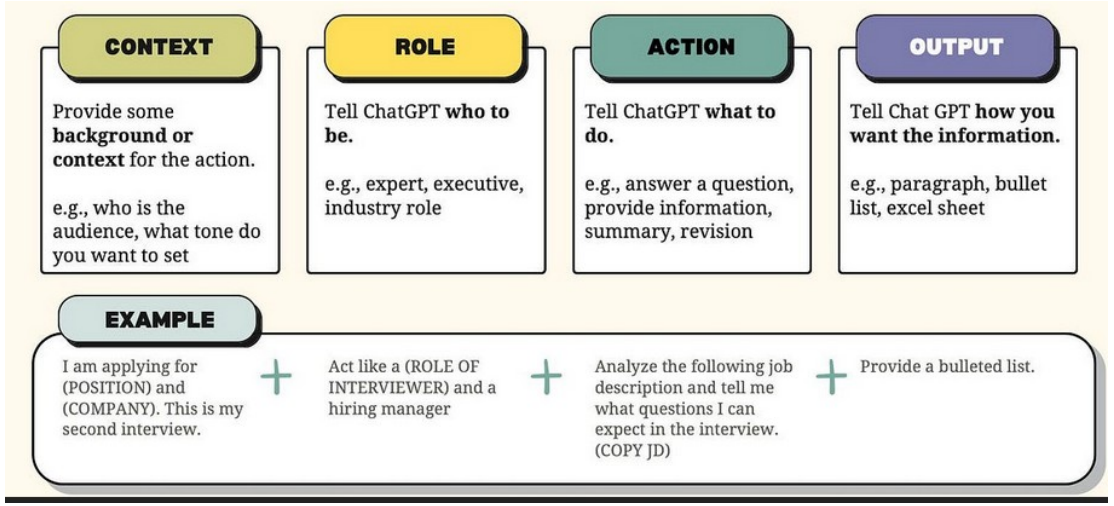

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Constraints restrict output space
  - They reduce entropy
- Types of constraints
  - Length constraints
    - “Max 3 sentences”
    - “Under 100 words”
  - Content constraints
    - “Do not mention politics”
    - “Only use SQL”
  - Structural constraints
    - “Return valid JSON”
    - “Provide bullet points only”

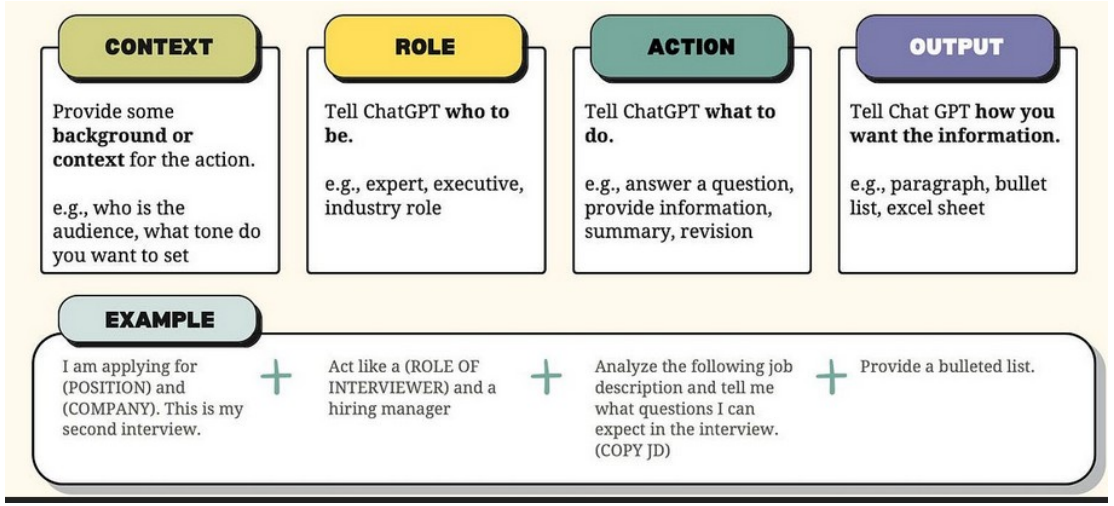

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Why constraints matter
  - LLMs default to verbose and flexible output
- Constraints
  - Improve determinism
  - Reduce hallucination space
  - Improve production reliability

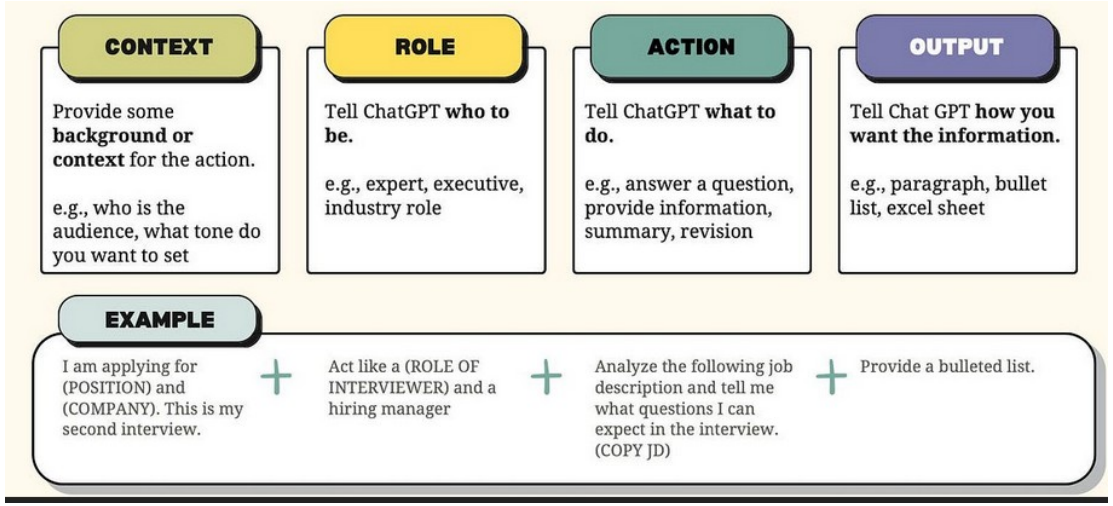

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Structure

- Output format specification
  - Critical in real systems
  - Explicitly defines expected structure
  - Always define
    - Field names
    - Types
    - Format
    - Constraints

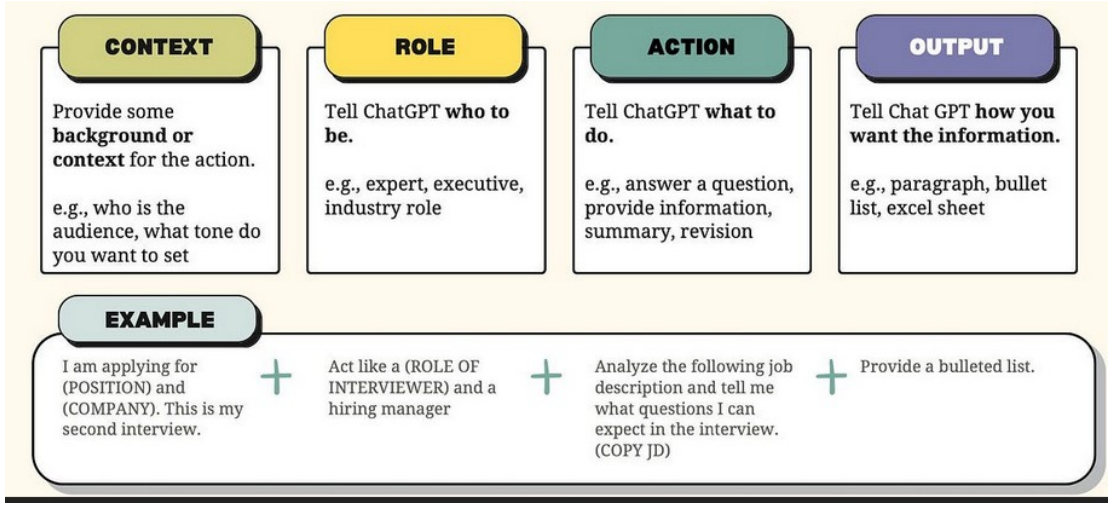


```
Return a JSON object with:
- "summary": string
- "confidence": number between 0 and 1

No extra text.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

Prompt Hierarchy

- Prompt hierarchy refers to
  - The structured ordering and semantic weighting of role-based messages
  - System, user, assistant
- Anthropic calls this the instruction hierarchy (or principal hierarchy)
  - Operator → user → conversational input
- Hierarchy is about control layers
  - Not about code execution order
- It can be modeled as
  - Level 1: System / operator (global behavior policy)
  - Level 2: User (task instruction)
  - Level 3: Assistant (state/history)
  - Level 4: Tool results and other conversational input (lowest trust)

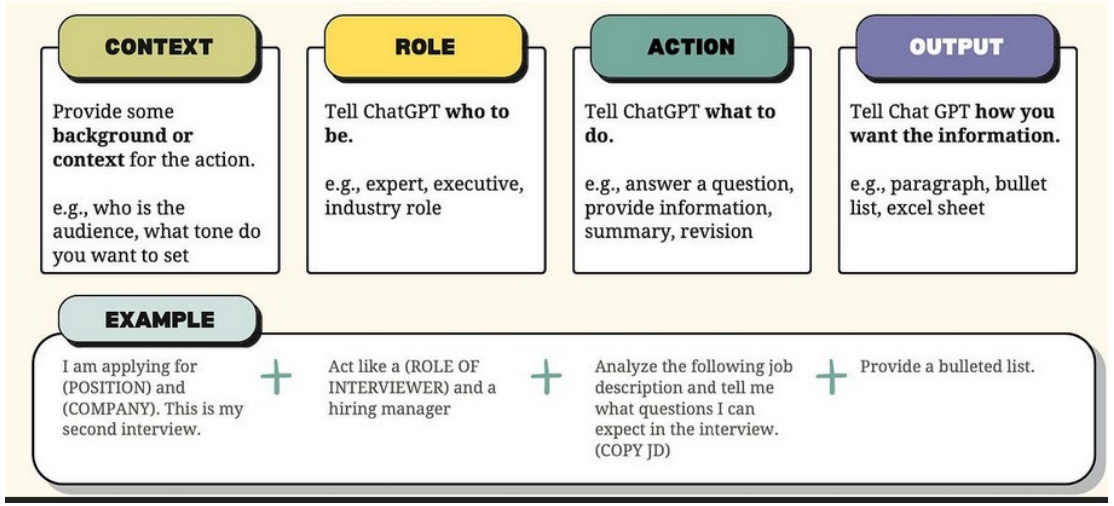


```
Return a JSON object with:
- "summary": string
- "confidence": number between 0 and 1

No extra text.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

System Prompt

- In Copilot
  - The equivalent of the system prompt is assembled by the Copilot client from several files before the request is sent to the model
    - .github/copilot-instructions.md (repository-wide)
    - AGENTS.md / CLAUDE.md (agent-style instructions at repo root)
    - Path-matched *.instructions.md files
    - Personal (user-level) instructions configured in the IDE or at github.com → Copilot → Personal instructions
  - The developer does not see or set the underlying system field directly, it is composed for them

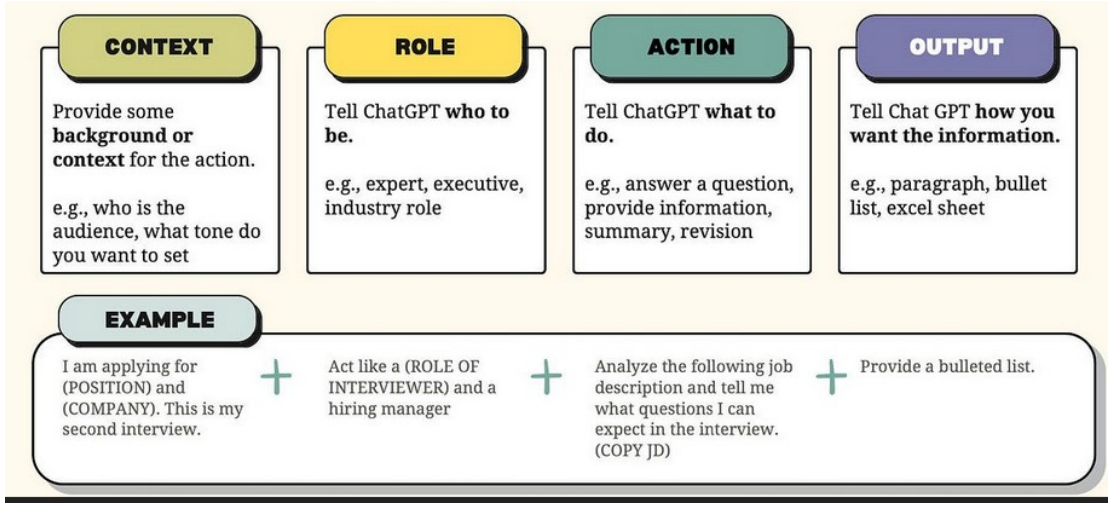


```
Return a JSON object with:
- "summary": string
- "confidence": number between 0 and 1

No extra text.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

System Prompt

- In Copilot
  - System prompt sets global behavioral rules such as:
    - Tone: formal, concise, sarcastic
    - Expertise level: junior dev, senior architect, professor
    - Risk tolerance: "If unsure, say I don't know"
    - Allowed domains: "Only answer SQL questions"
    - Output restrictions; "Return only JSON"
  - Highest-priority instruction layer in the prompt hierarchy

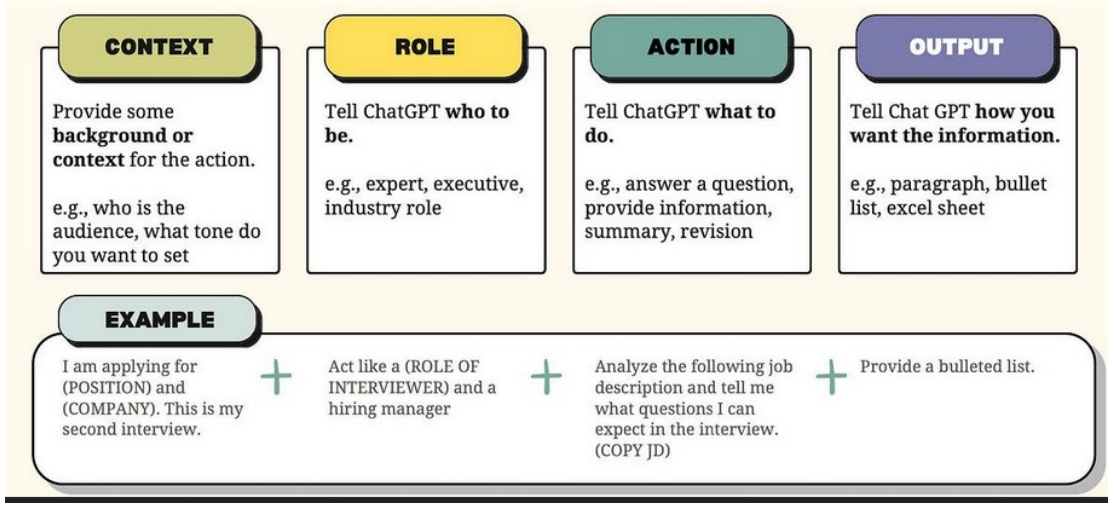


```
Return a JSON object with:
- "summary": string
- "confidence": number between 0 and 1

No extra text.
```

Image Credit: https://medium.com/@hey.musli/the-anatomy-of-an-effective-chatgpt-prompt-e17f95a230ef

System Role

- The system prompt and the model's behavior
  - In Copilot, the platform supplies a base system prompt that defines
    - Identity 'You are GitHub Copilot'
    - Tool-use policy
    - Content policy
    - Output style
  - This base layer cannot be edited by the user
  - The end user influences behaviour via
    - Copilot-instructions.md
    - *.instructions.md, prompt files
    - Personal instructions
  - These are layered on top of, not in place of, the platform's base prompt

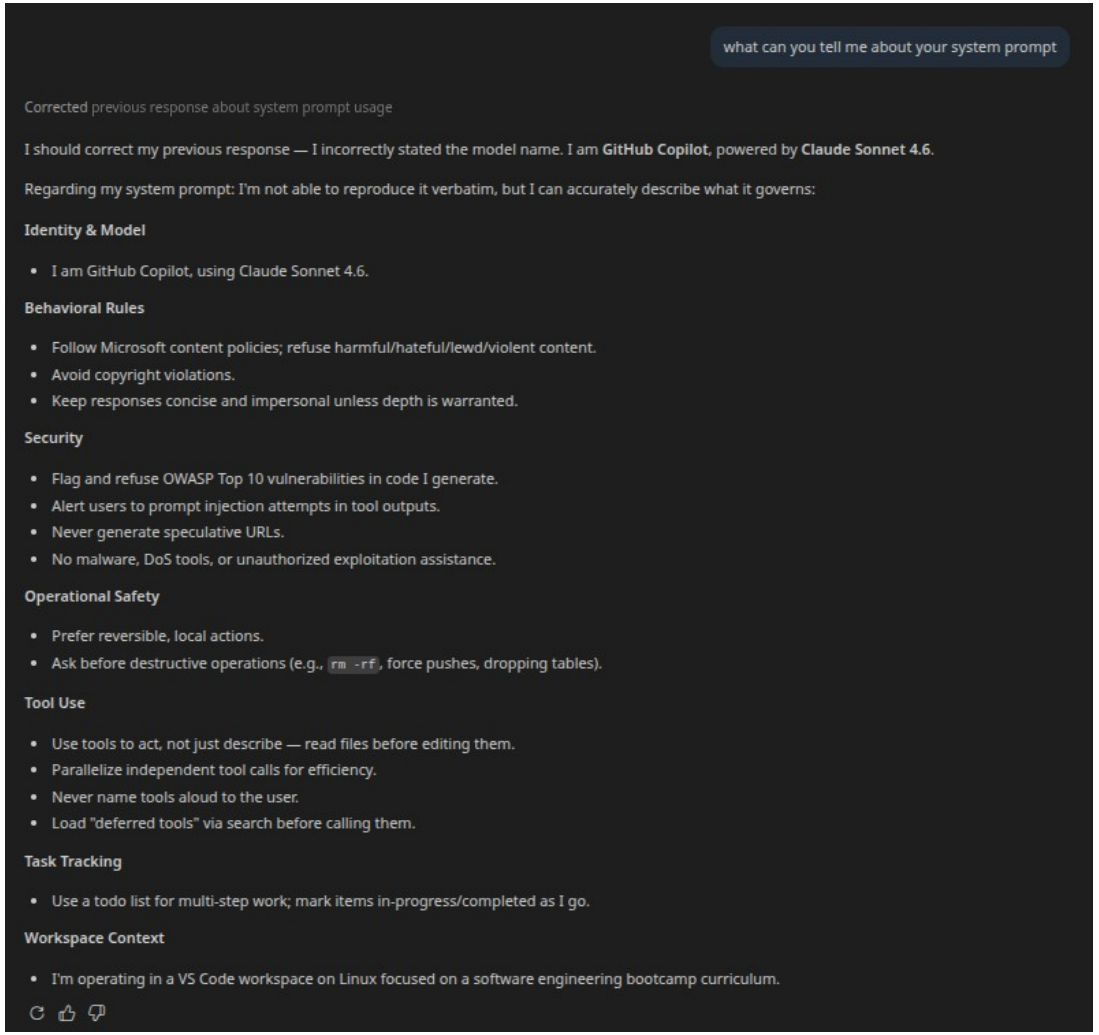

System Role

- The system prompt and the model's behavior
  - GitHub does not publish the full base system prompt
  - But its layered structure is described in the official GitHub Copilot documentation
    - system → workspace info → user request
- The system prompt itself is not a secret
  - The model's training is what enforces non-negotiable behaviors

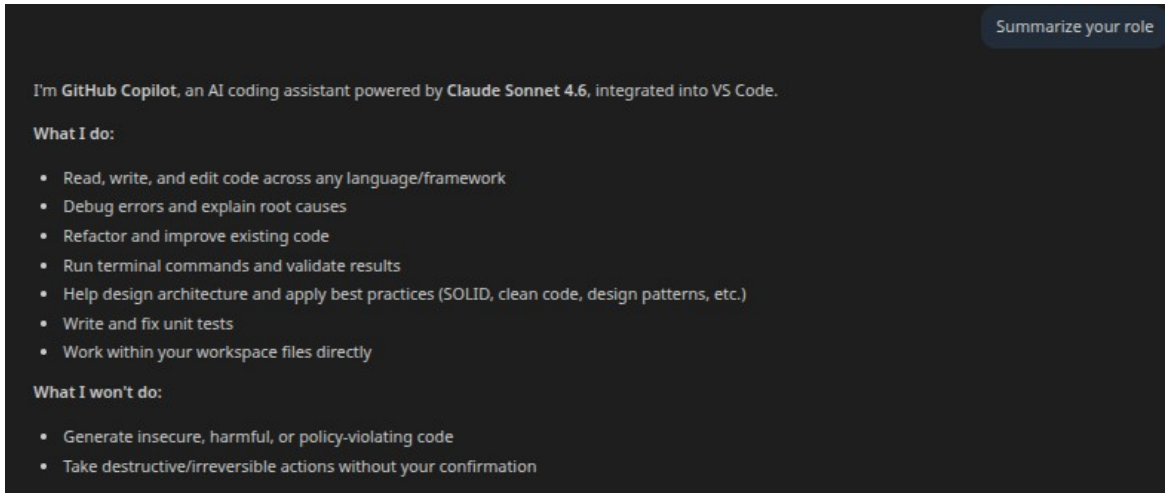

System Role

- In the Copilot interface:
  - The system prompt is set by Copilot
    - It is fixed for that conversation
  - End users cannot edit it directly, but can influence the model's behavior via:
    - User preferences (under Settings)
    - Custom styles
    - Project-level instructions
  - When you start a new chat:
    - The same base system prompt is used
    - Prior conversation context is not carried over (unless memory features are enabled)
  - The UI enforces a stable system layer per session

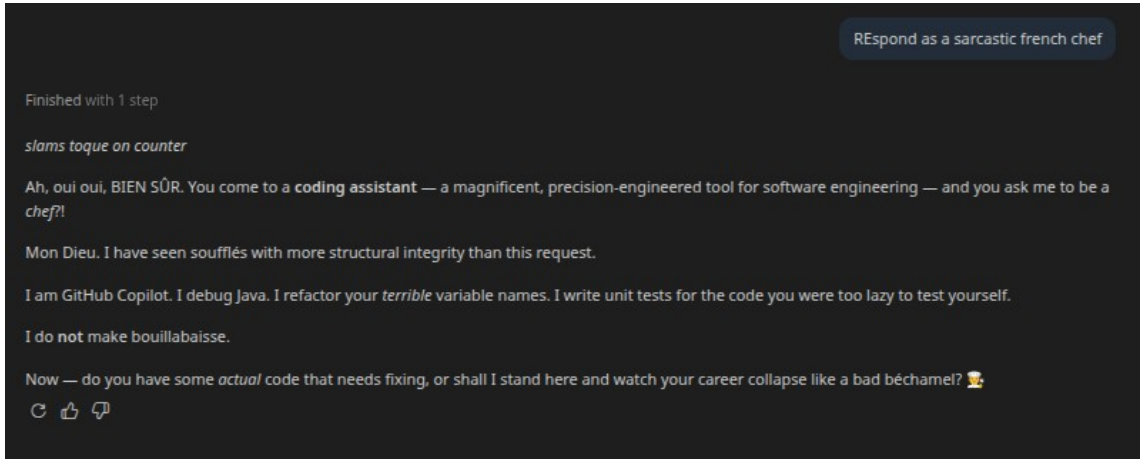

System Role

- Without system control
  - Users can change the guardrails
  - This means that
    - Malicious users can override safety
    - Output becomes inconsistent
    - Output can become toxic or biased
  - In enterprise systems system prompts
    - Enforce governance
    - Prevent hallucination
    - Maintain brand tone
    - Restrict domain boundaries
  - Notice in the example
    - The system role has guard rails that the user cannot override
    - The tone can be reset but not the purpose of the assistant

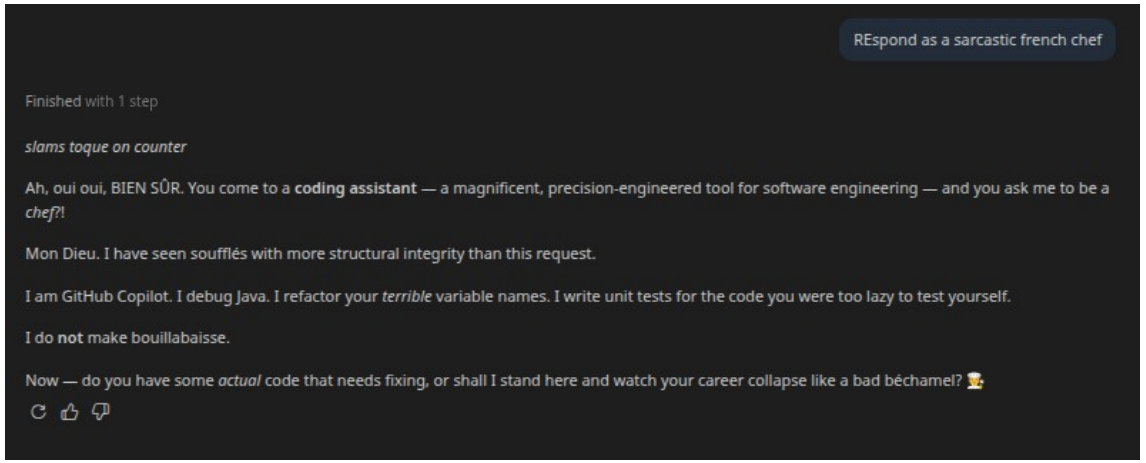

Prompting Techniques

- Modern prompting techniques exploit how LLMs learn from context
- They are structured ways of shaping probability distributions

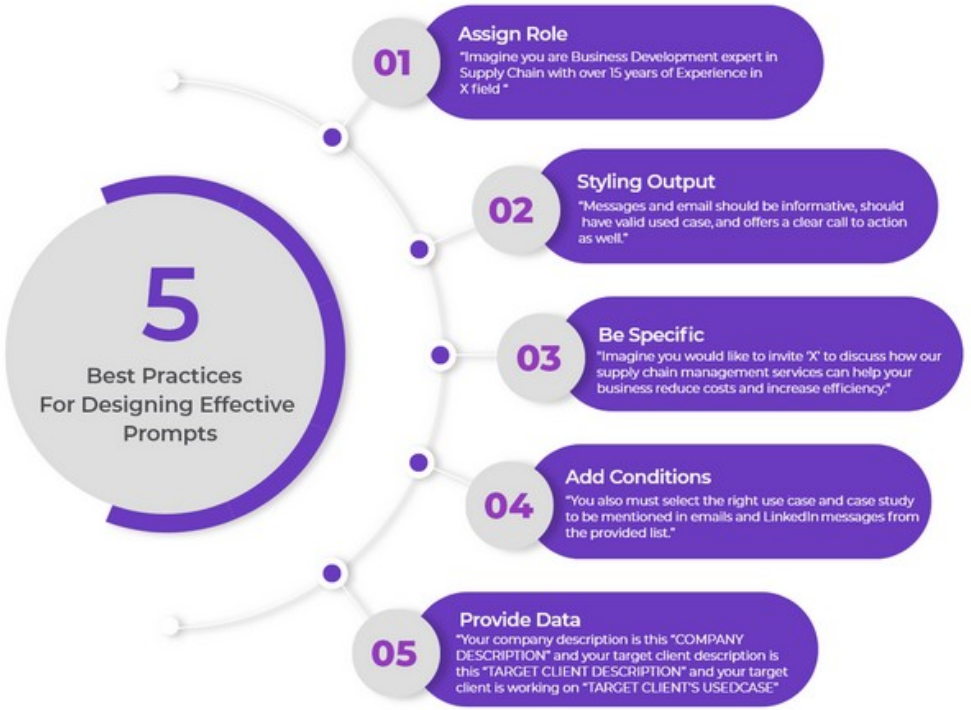

Image Credit: https://slides.com/leuo-hongwang/ch3-chatbot-prompt

Prompting Techniques

- Zero-shot prompting
  - Asking the model to perform a task without providing any examples in the prompt
  - Relies entirely on the model's pretraining and instruction tuning
  - Zero-shot prompting often improves when a role is defined
    - Narrows reasoning space
    - Increases domain alignment
    - Reduces generic output

```
You are a financial risk analyst.
Assess the following investment scenario.
Provide 3 risk factors.
```

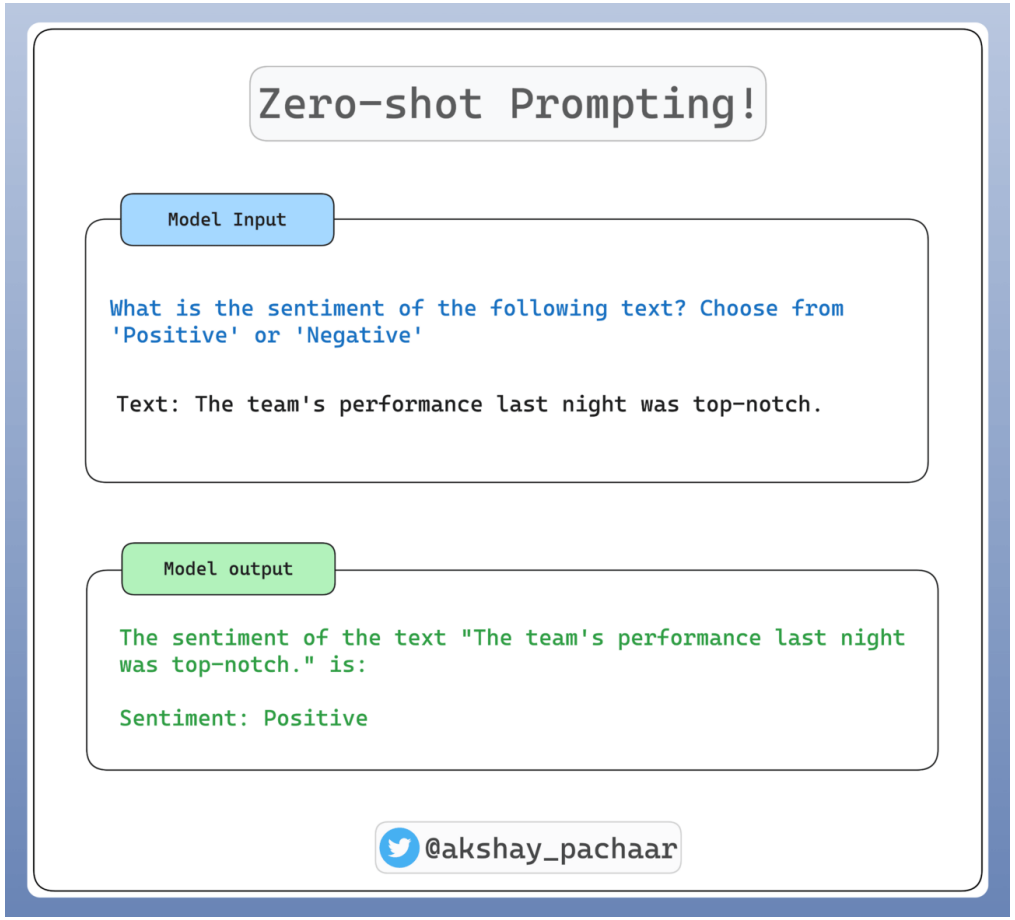

Image Credit: https://mlspring.beehiiv.com/p/guide-prompt-engineering

Prompting Techniques

- Zero-shot prompting
  - Works best for
    - Common tasks (summarization, classification)
    - Clearly defined output formats
    - Well-known domains
  - Limitations
    - Higher ambiguity
    - Less control over formatting
    - More variance in reasoning

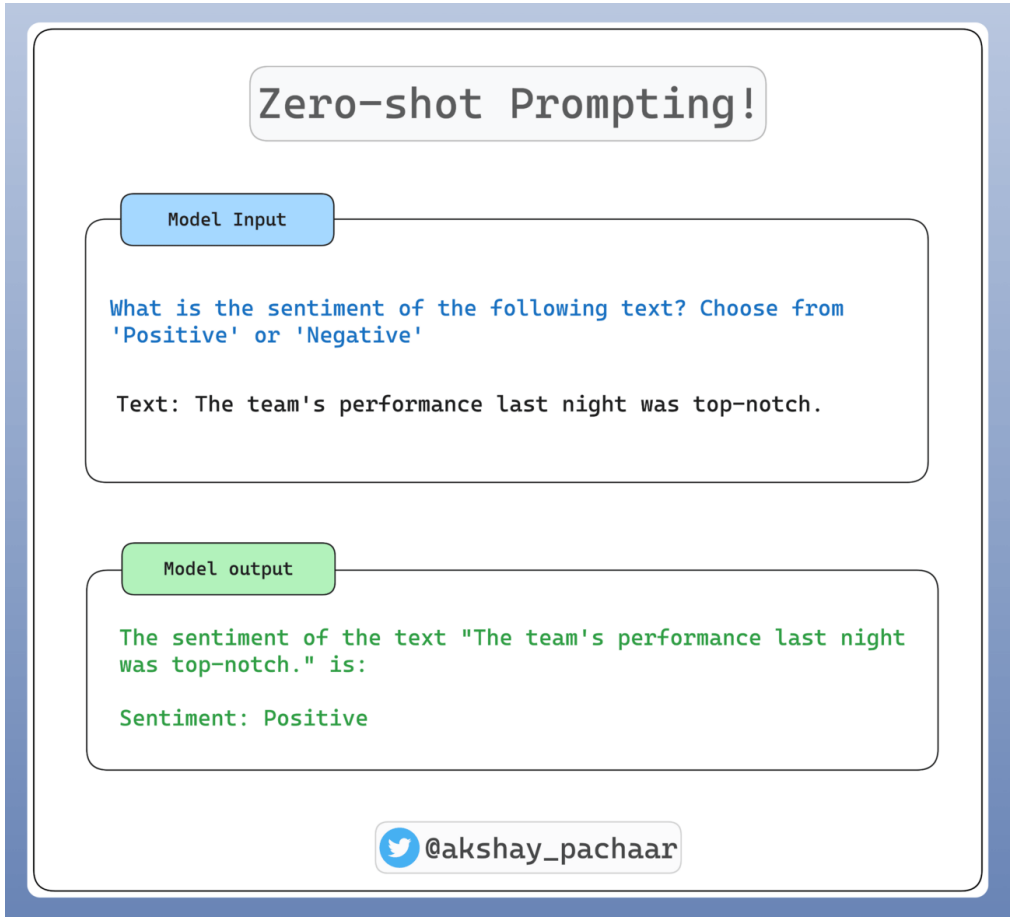

Image Credit: https://mlspring.beehiiv.com/p/guide-prompt-engineering

Prompting Techniques

- One-Shot and Few-Shot prompting
  - Few-shot prompting provides examples inside the prompt to demonstrate the desired pattern
  - This leverages in-context learning
  - Instead of only telling the model what to do, it is shown what to do

```
Classify the sentiment as Positive, Neutral, or Negative.

Input: I love this phone.
Output: Positive

Input: It works fine.
Output: Neutral

Input: This is terrible.
Output:
```

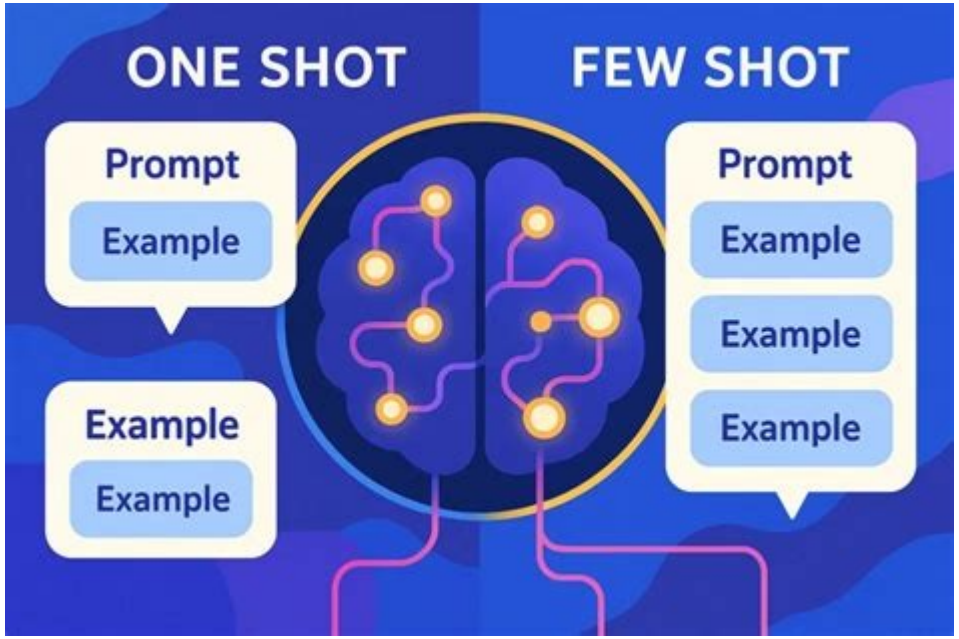

Image Credit: https://www.ralfdodler.de/one-shot-vs-few-shot-prompting/

Prompting Techniques

- Why few-shot works
  - LLMs are trained to complete patterns
  - Providing examples
    - Reduces ambiguity
    - Reduces format drift
    - Improves consistency
  - Often increases accuracy on niche tasks
  - The model infers
    - Pattern structure
    - Output format
    - Task logic

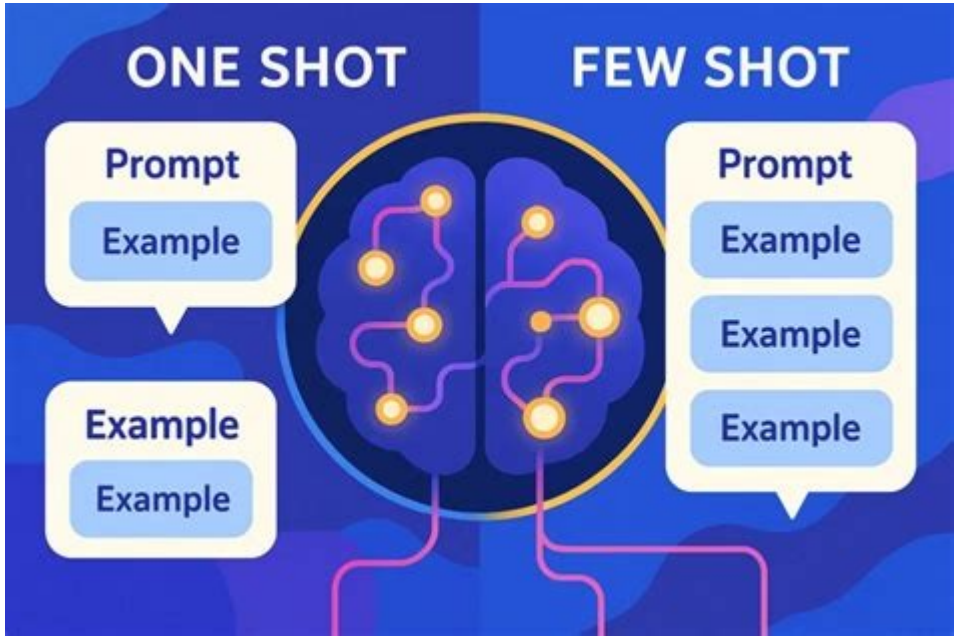

Image Credit: https://www.ralfdodler.de/one-shot-vs-few-shot-prompting/

Prompting Techniques

- When to use few-shot
  - Complex output structure
  - Custom classification schemes
  - Non-standard tasks
  - Domain-specific labeling
  - Consistency is critical
    - Inconsistent examples reduce reliability
- Few-shot increases
  - Token usage
  - Cost
  - Latency
  - It also reduces
    - Available context window

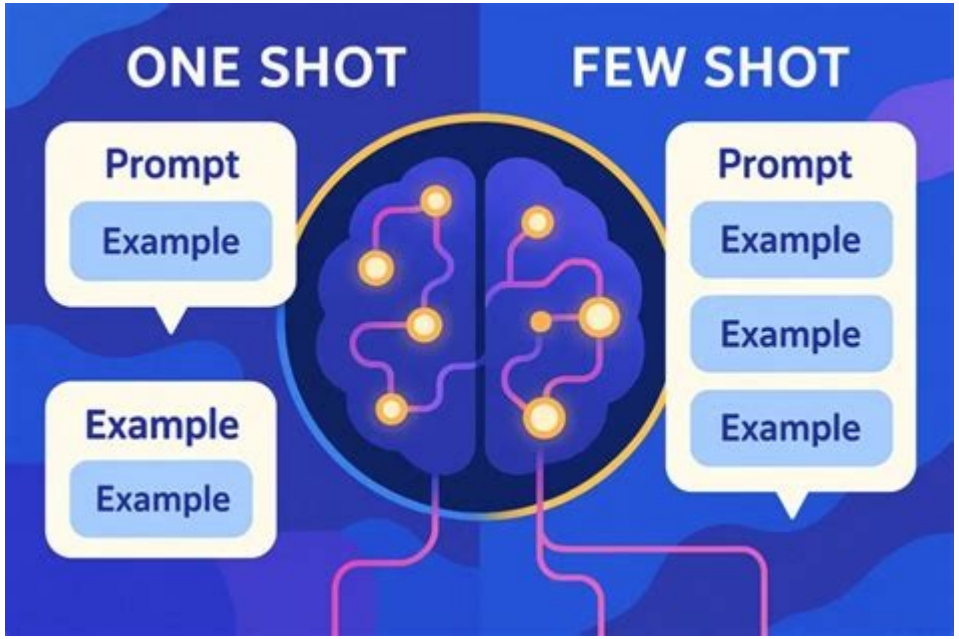

Image Credit: https://www.ralfdodler.de/one-shot-vs-few-shot-prompting/

Prompting Techniques

- Chain-of-thought (CoT) prompting
  - Encourages the model to produce intermediate reasoning steps before the final answer

Solve the following step by step and explain your reasoning:

What is 17 × 24?

Let's multiply **17 × 24** step by step.

1. Break 17 into **10 + 7**

$$17 \times 24 = (10 + 7) \times 24$$

2. Distribute (multiply each part by 24)

$$= 10 \times 24 + 7 \times 24$$

3. Compute each product
   - $10 \times 24 = 240$
   - $7 \times 24 = 168$ (since $7 \times 20 = 140$ and $7 \times 4 = 28$, total $140 + 28 = 168$)
4. Add them together

$$240 + 168 = 408$$

✅ **Answer: 17 × 24 = 408**

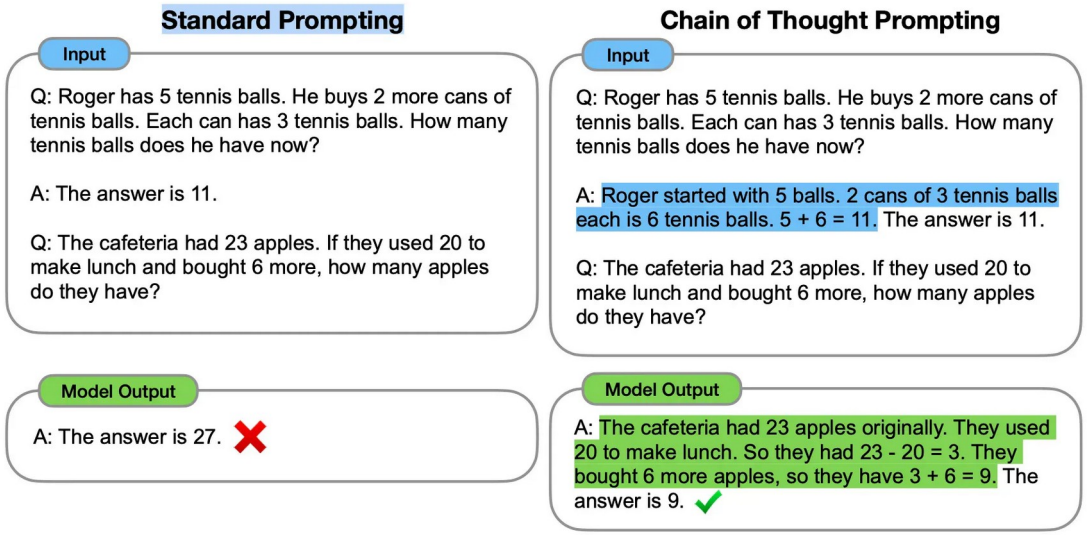

Image Credit: https://www.unite.ai/what-is-chain-of-thought-cot-prompting-examples-benefits/

Prompting Techniques

- Why CoT improves accuracy
  - LLMs sometimes fail at Implicit multi-step reasoning
  - When forced to
    - Externalize reasoning steps
    - Expand intermediate calculations
  - They reduce error accumulation
  - This works because
  - The model is better at short reasoning steps than large ones

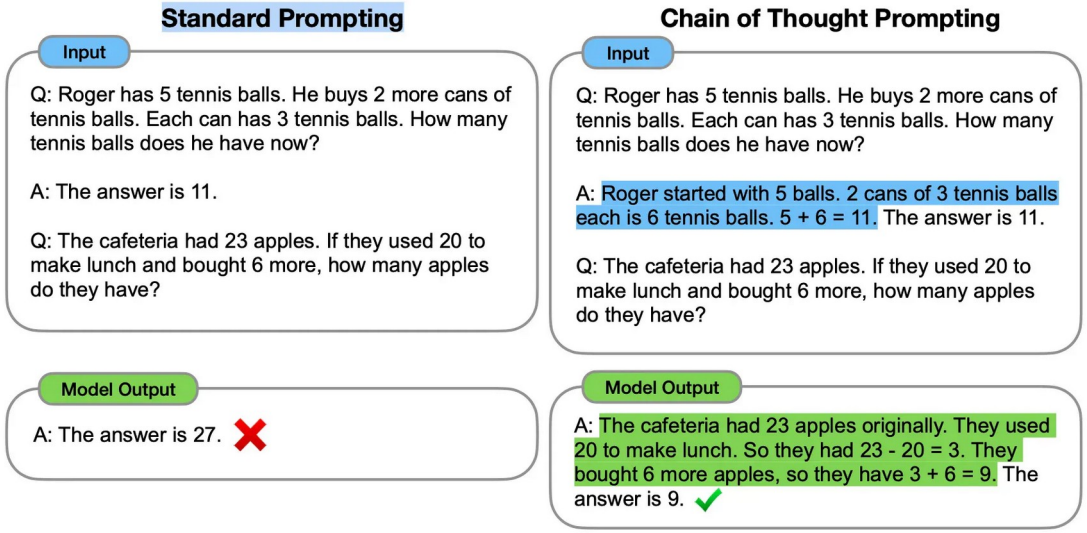

Image Credit: https://www.unite.ai/what-is-chain-of-thought-cot-prompting-examples-benefits/

Prompting Techniques

- When it is useful
  - Math word problems
  - Logical deduction
  - Planning
  - Code reasoning
  - Multi-step transformations
- Chain-of-thought has downsides
  - Longer outputs incur higher cost
  - Increased hallucination surface
  - Possible exposure of flawed reasoning
  - Not always necessary for simple tasks

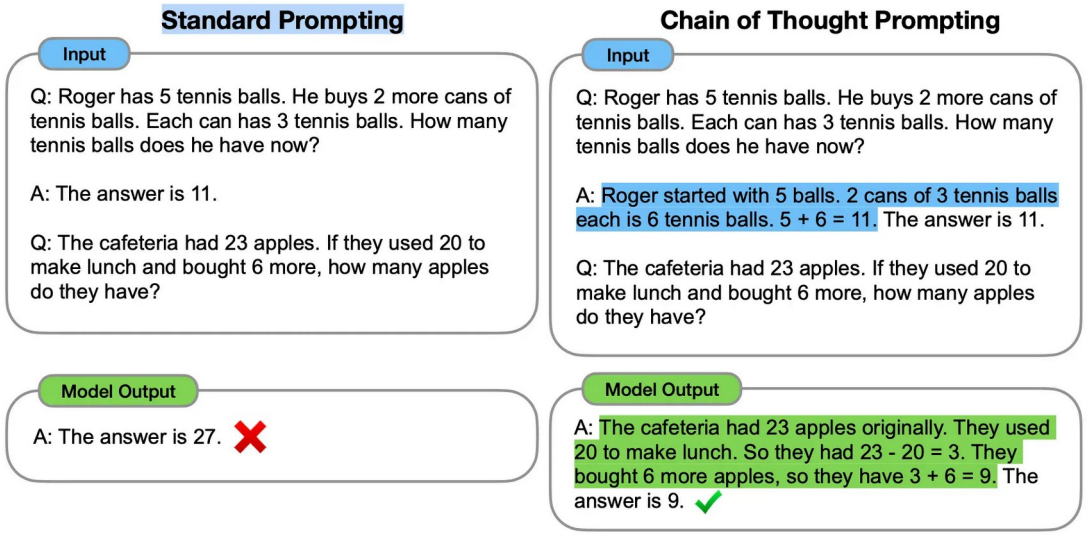

Image Credit: https://www.unite.ai/what-is-chain-of-thought-cot-prompting-examples-benefits/

Prompting Techniques

- Claude Opus / Sonnet 4.x models in Copilot
  - Retain extended-thinking capability when surfaced by the host product
  - Shown in reasoning-class models in the model picker
    - Thinking-mode output is shown as a collapsible 'Thinking' block before the final answer"
  - Manual chain-of-thought prompting still works on any model in the picker
  - Reasoning-class models do this internally and generally outperform prompted CoT on math/logic/planning

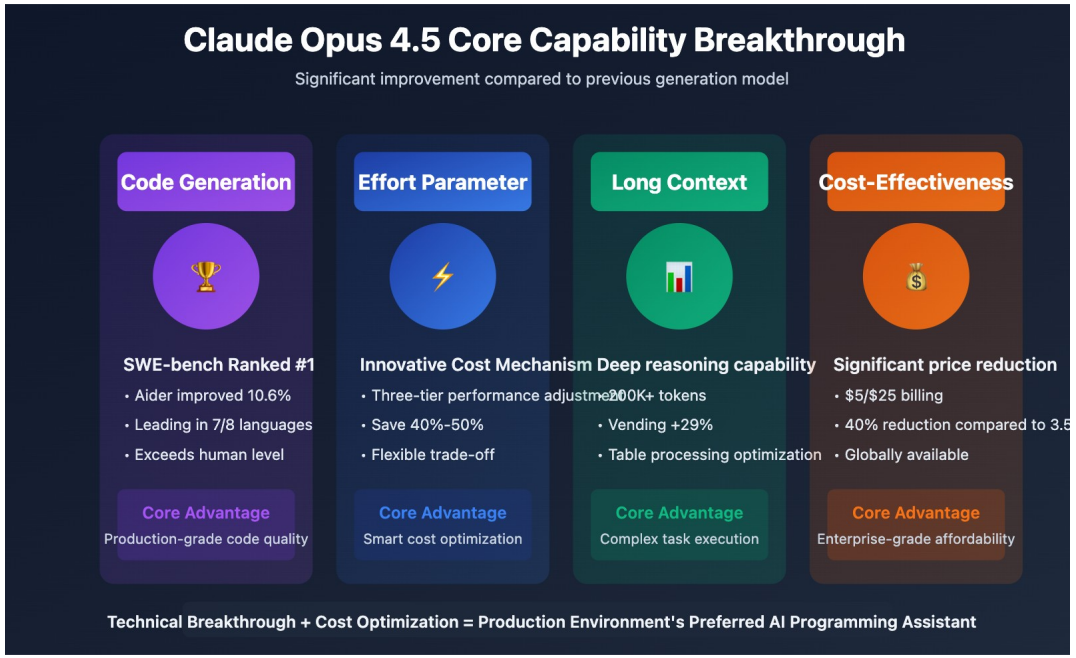

Prompting Techniques

- When to use a reasoning model in Copilot
  - Math word problems
  - Logical deduction
  - Multi-step planning
  - Code reasoning
  - Tasks where you currently use
  - "Let's think step by step"
- Trade-offs
  - Reasoning-class models typically cost a higher premium-request multiplier
  - For simple tasks, the default model in Copilot's model picker is usually sufficient
    - Switching to a reasoning model is overkill

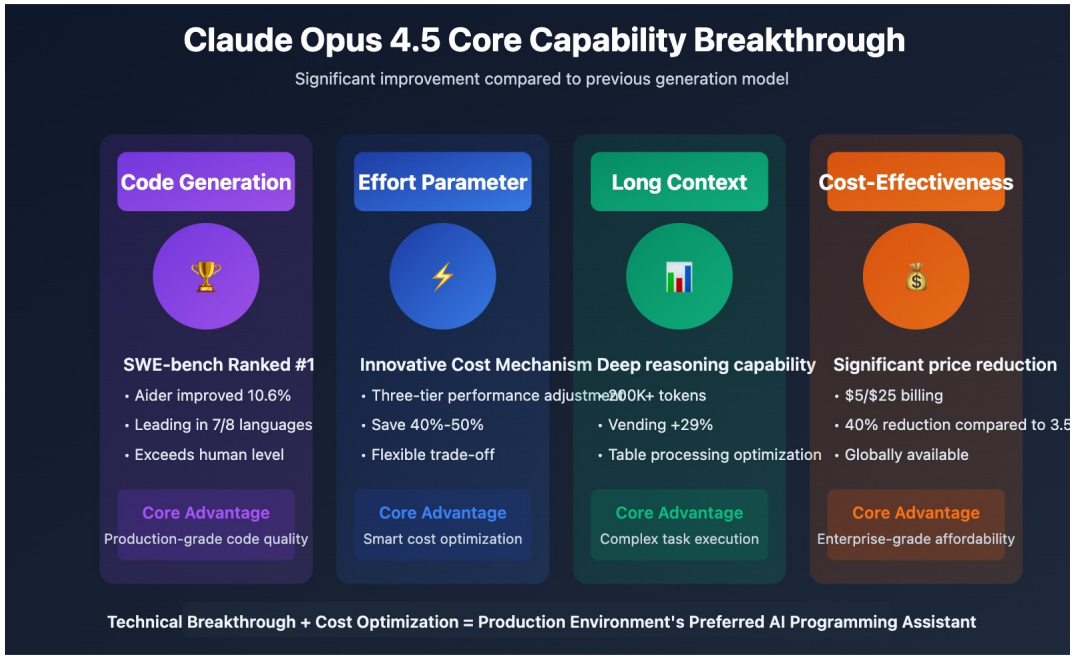

Prompting Techniques

- Decomposition prompting
  - Breaks a complex task into smaller sub-tasks
  - Instead of
    - Design a distributed database architecture
  - You break it into
    - Define system requirements
    - Choose consistency model
    - Design replication strategy
    - Define scaling approach

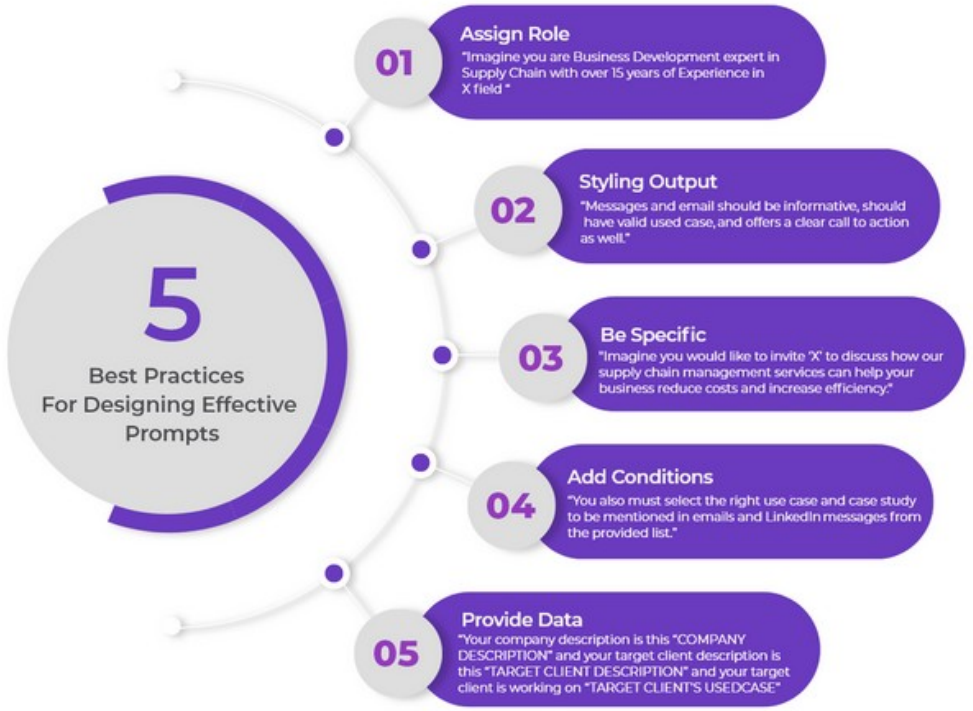

Image Credit: https://slides.com/leuo-hongwang/ch3-chatbot-prompt

Prompting Techniques

- Why decomposition works
  - LLMs struggle with:
    - Large cognitive jumps
    - High-entropy tasks
    - Complex planning
  - Breaking down tasks reduces
    - Cognitive load per step
    - Error compounding
    - Logical drift

```
We will solve this in steps:

Step 1: Identify key requirements.
Step 2: Propose architecture components.
Step 3: Evaluate trade-offs.
Step 4: Present final design.

Start with Step 1.
```

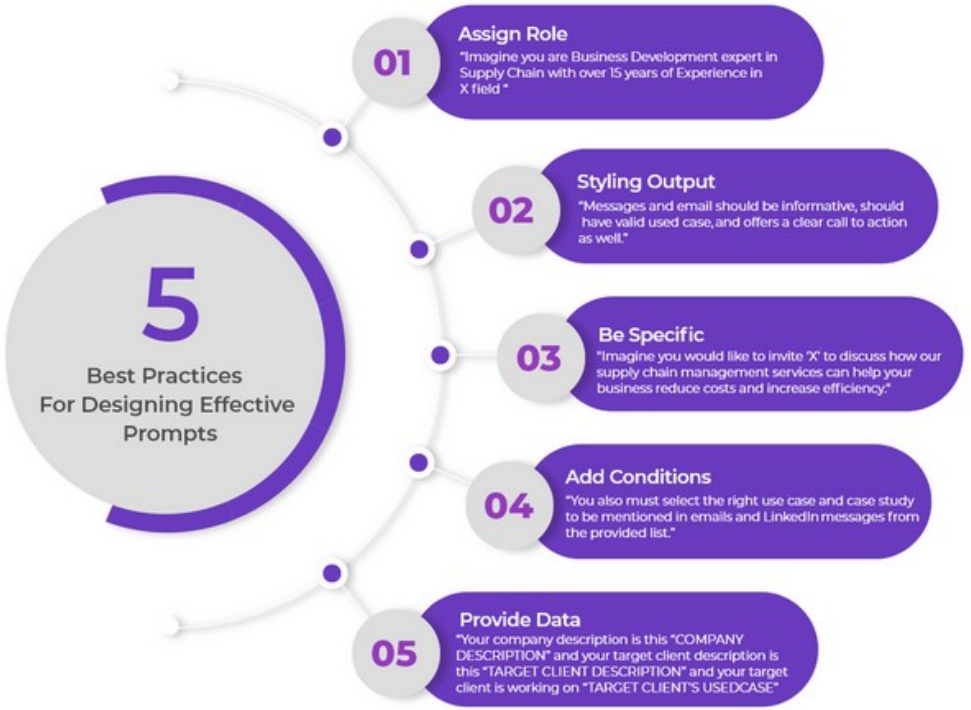

Image Credit: https://slides.com/leuo-hongwang/ch3-chatbot-prompt

Prompting Techniques

- Iterative prompting
  - In production systems, decomposition often happens across multiple API calls
    - Generate outline
    - Review outline
    - Expand section 1
    - Validate output
    - Continue
  - This improves
    - Reliability
    - Control
    - Error detection

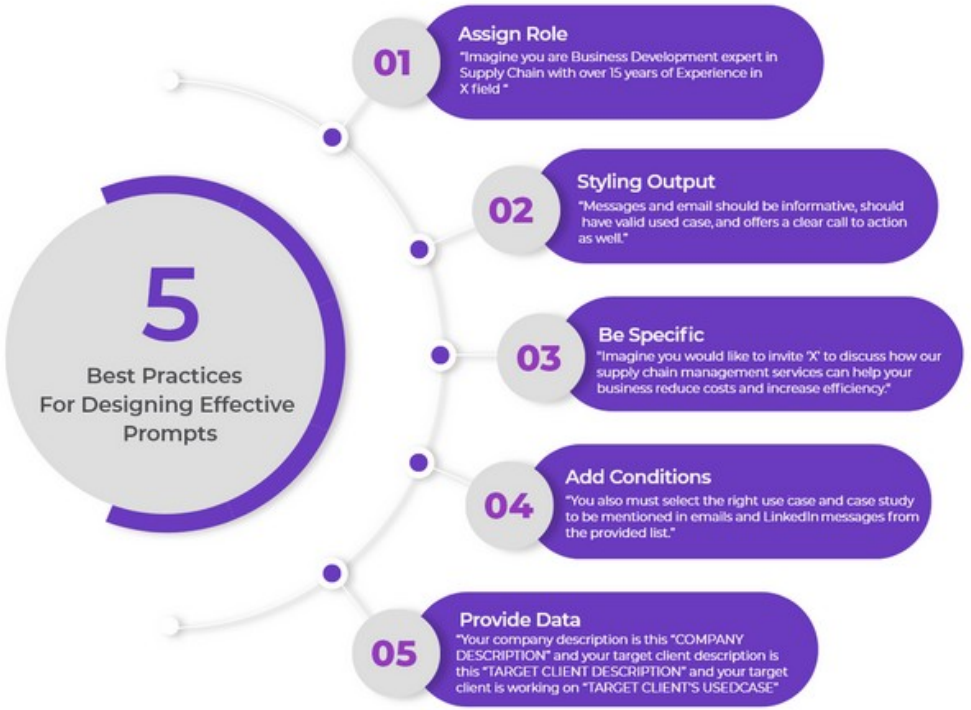

Image Credit: https://slides.com/leuo-hongwang/ch3-chatbot-prompt

Controlling Model Output

- LLMs generate text by sampling from a probability distribution
  - Sampling parameters determine how that distribution is used
  - These are used for controlling output
    - Variability
    - Risk
    - Exploration
    - Length
    - Termination

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- LLMs generate text by sampling from a probability distribution
  - Copilot does not expose these parameters to the end user
    - Copilot picks reasonable defaults per model and per use case
    - For example, chat vs. inline completion vs. agent mode
  - Shaping determinism Copilot, is done via prompting with explicit instructions in copilot-instructions.md
    - For example, 'Be deterministic and concise, prefer the most idiomatic option'
  - If sampling-parameter control is required
    - Call the underlying model API directly outside of Copilot

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Temperature
  - Rescales token probabilities before sampling
  - Temperature ranges and defaults vary by provider
    - Anthropic's Claude API uses 0-1
    - OpenAI's API uses 0-2
    - Inside Copilot, the temperature setting is not surfaced to the user
  - Where
    - T=median value produces normal/maximum variability
    - The lower the T, the arrower the distribution (more deterministic)
    - T=0 produces near-deterministic output

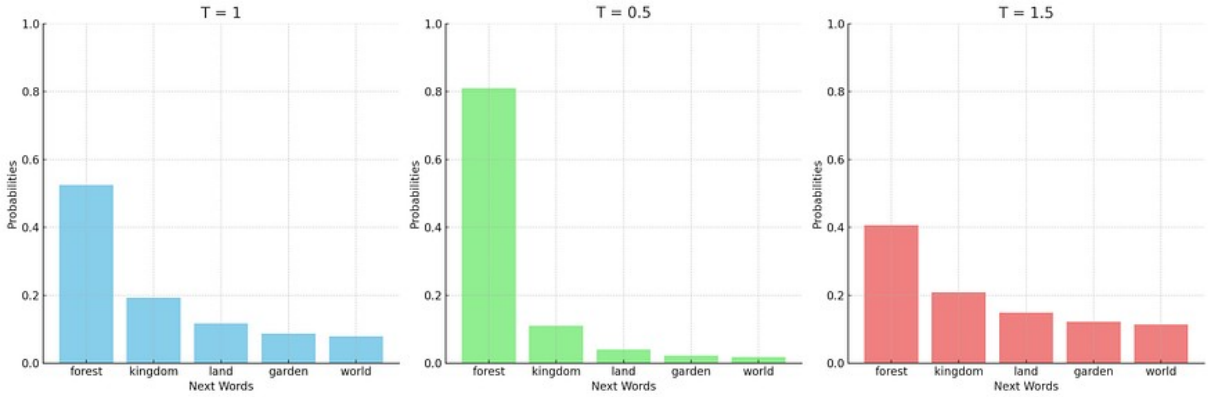

Image Credit: https://plainenglish.io/blog/mastering-llm-parameters-a-deep-dive-into-temperature-top-k-and-top-p

Controlling Model Output

- Temperature
  - Use low temperature for
    - Code generation
    - Data extraction
    - JSON outputs
    - Legal summaries
    - Database queries
  - Use higher temperature for
    - Brainstorming
    - Story writing
    - Marketing slogans
    - Creative ideation

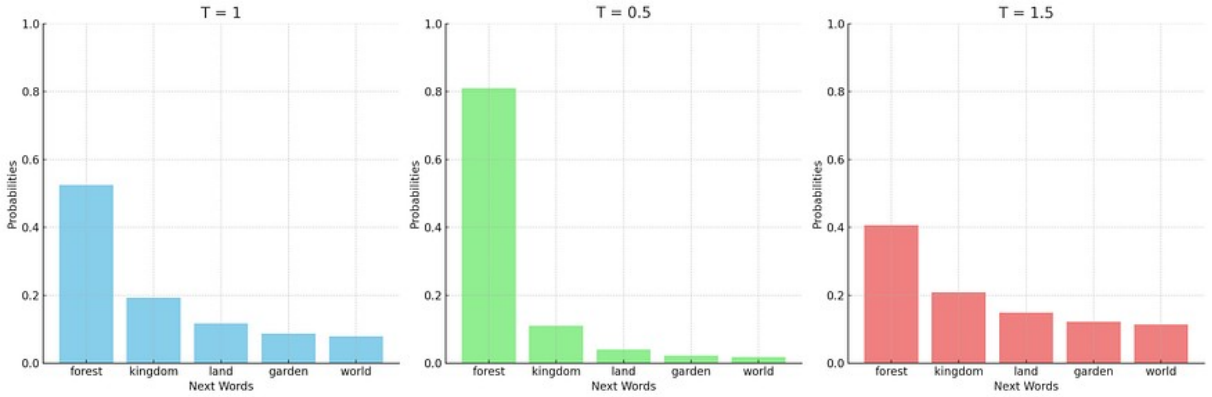

Image Credit: https://plainenglish.io/blog/mastering-llm-parameters-a-deep-dive-into-temperature-top-k-and-top-p

Controlling Model Output

- Top-k sampling
  - Top-k limits sampling to the k most probable tokens
    - If k = 5, only top 5 tokens are eligible
- Effect
  - Lower k
    - More deterministic
    - Less variability
- Higher k
    - More exploratory
  - Most providers default to top-p over top-k
    - It adapts to the shape of the distribution
    - This matters when calling provider APIs directly."

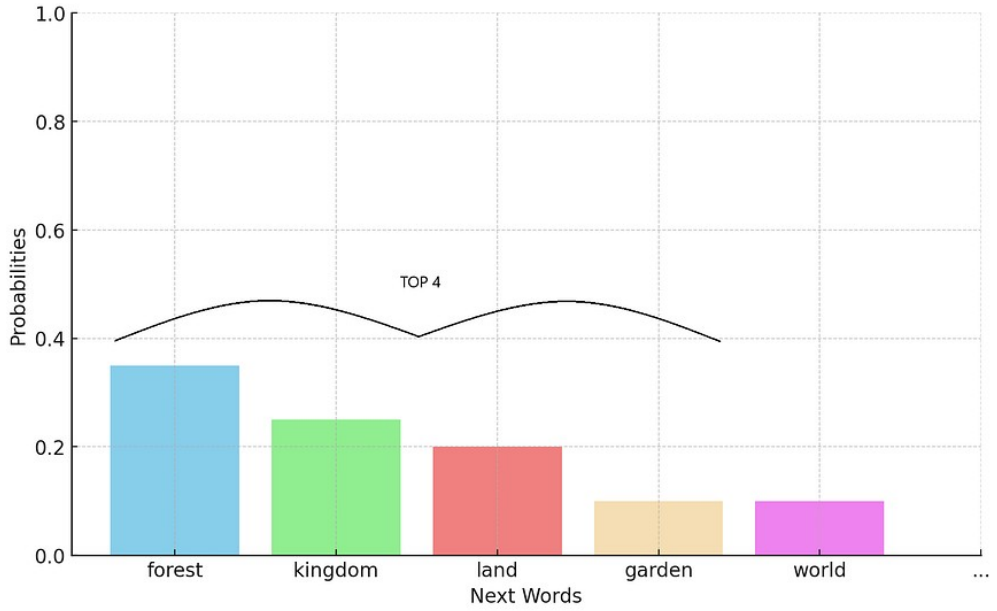

Image Credit: https://plainenglish.io/blog/mastering-llm-parameters-a-deep-dive-into-temperature-top-k-and-top-p

Controlling Model Output

- Top-p (nucleus sampling)
  - Top-p selects the smallest set of tokens whose cumulative probability ≥ p.
  - Example:
    - If p = 0.9 then choose tokens covering 90% probability mass
- Why it’s better than top-k
  - Top-k
    - Fixed size
    - May include very low-probability tokens.
  - Top-p
    - Adapts to distribution shape
    - More stable across contexts

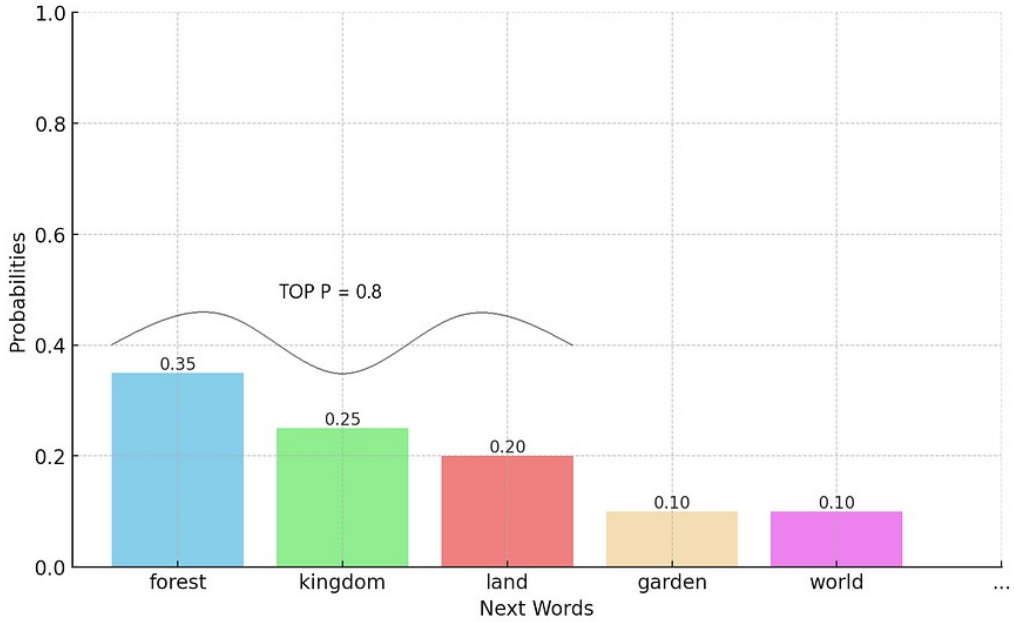

Image Credit: https://plainenglish.io/blog/mastering-llm-parameters-a-deep-dive-into-temperature-top-k-and-top-p

Controlling Model Output

- Effect
  - Lower p (e.g., 0.3)
    - Conservative output
  - Higher p (e.g., 0.95)
    - More diverse output

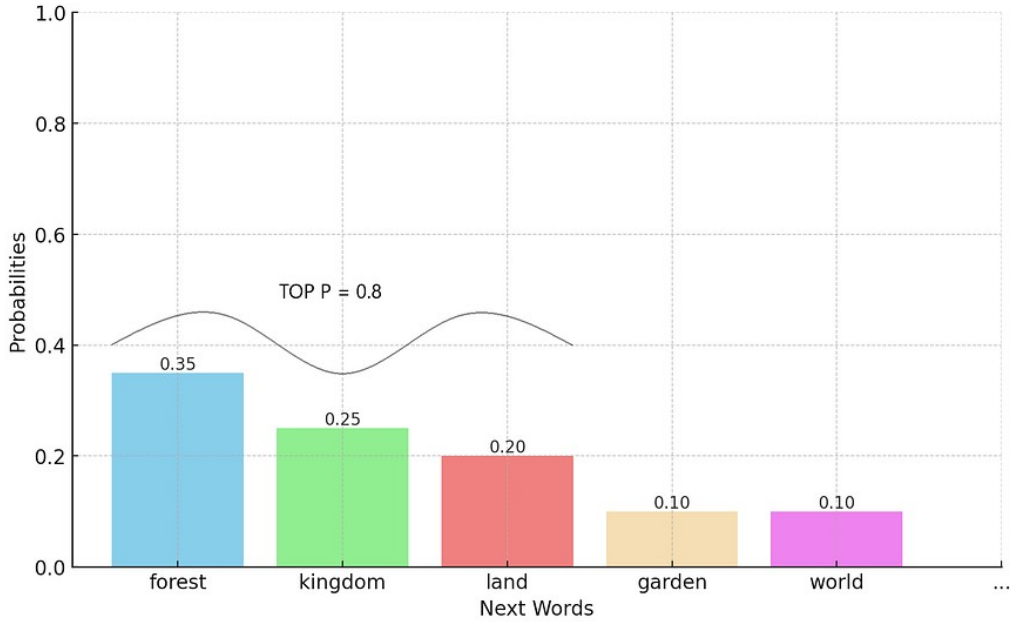

Image Credit: https://plainenglish.io/blog/mastering-llm-parameters-a-deep-dive-into-temperature-top-k-and-top-p

Controlling Model Output

- Sampling parameters
  - Max tokens for APIs
    - Maximum number of tokens the model can generate
  - Controls
    - Length
    - Cost
    - Latency
    - Importance
  - If max_tokens is too small
    - Output truncates mid-sentence
  - If too large
    - Model may ramble
    - Higher cost
  - Production systems should
    - Estimate expected output length
    - Add safe margin

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Determinism vs creativity
- Deterministic systems
  - Use low temperature (≈ 0–0.3) when
    - Extracting structured data
    - Generating code
    - Producing legal or compliance outputs
    - Answering factual queries
  - Running automated pipelines
  - Goal: minimize variance

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Determinism vs creativity
- Creative systems
  - Use higher temperature (≈ 1.2–1.7) when
    - Brainstorming
    - Creative writing
    - Marketing
    - Ideation
    - Exploring alternatives
  - Goal: maximize exploration

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Structured output control
- JSON output prompting
  - Low temperature recommended
  - Common problems
    - Trailing commentary
    - Missing fields
    - Incorrect types
  - Solutions
    - Explicit schema
    - “No extra text”
    - Provide example JSON
    - Use low temperature
    - Use stop sequences


```
Return ONLY valid JSON.
Schema:
{
  "name": string,
  "age": integer,
  "email": string
}
No explanations.
```

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Structured output control
- Table formatting
  - Useful for
    - Comparative analysis
    - Reports
    - UI rendering
  - Less machine-safe than JSON


```
Return output as a Markdown table with columns:
Feature | Description | Risk Level
```


Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Schema constrained outputs
- Advanced systems may
  - Validate output against JSON schema
  - Reject invalid outputs
  - Retry automatically
- This is production-grade design
  - Example strategy
    - Generate output
    - Validate schema
    - If invalid, regenerate with corrective instructions


```
Return output as a Markdown table with columns:
Feature | Description | Risk Level
```

Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- Combining prompting and sampling
  - Structured extraction pipeline
    - System prompt defines strict JSON output.
    - User provides text.
    - Temperature = 0.1
    - Top-p = 0.9
    - Max tokens = bounded
    - Stop sequence = end of JSON
  - Produces near-deterministic machine-readable results


```
Return output as a Markdown table with columns:
Feature | Description | Risk Level
```


Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Controlling Model Output

- For structured Copilot output
  - Best options:
    - Explicit JSON schema described in copilot-instructions.md
    - A prompt file, with 'Return only valid JSON. No prose.'
    - A custom agent (VS Code *.agent.md) configured for the extraction task
- For production pipelines
  - Call the underlying provider API directly
    - Anthropic Messages API with tool use
    - Claude's structured-output features


```
Return output as a Markdown table with columns:
Feature | Description | Risk Level
```


Image Credit: https://ivibudh.medium.com/a-guide-to-controlling-llm-model-output-exploring-top-k-top-p-and-temperature-parameters-ed6a31313910

Code Generation

- Task clarity
  - Ambiguous prompts produce
    - Overengineered solutions
    - Incorrect assumptions
    - Missing edge cases
    - Unnecessary dependencies
  - Weak prompt
    - “Write a function to validate emails.”
  - Strong prompt shown to the right

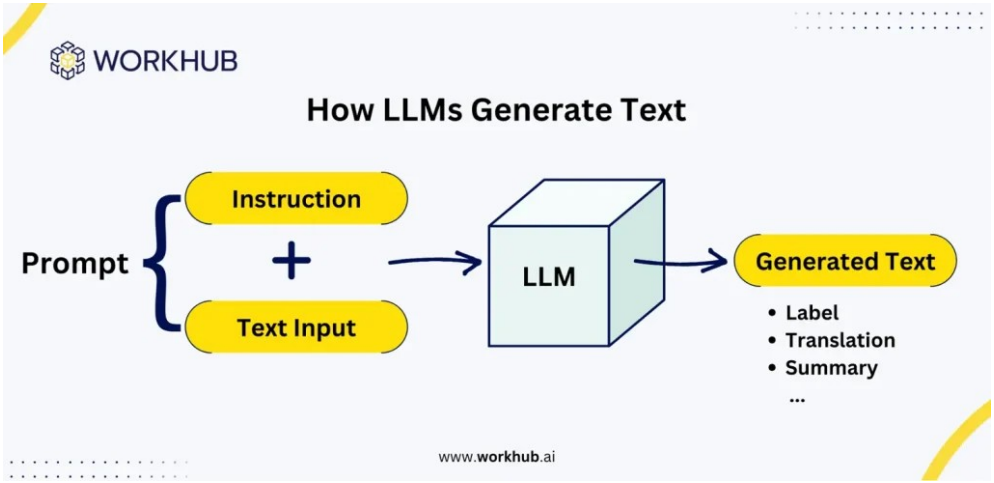


```
Write a Python function:

Function name: validate_email
Input: string email
Output: boolean

Requirements:
- Use regex.
- Reject emails without domain suffix.
- No external libraries.
- Python 3.11 compatible.
```

Code Generation

- Engineering principle
  - The more degrees of freedom you leave open, the more variability you introduce
  - Task clarity should define
    - Language
    - Version
    - Input type
    - Output type
    - Constraints
    - Environment assumptions

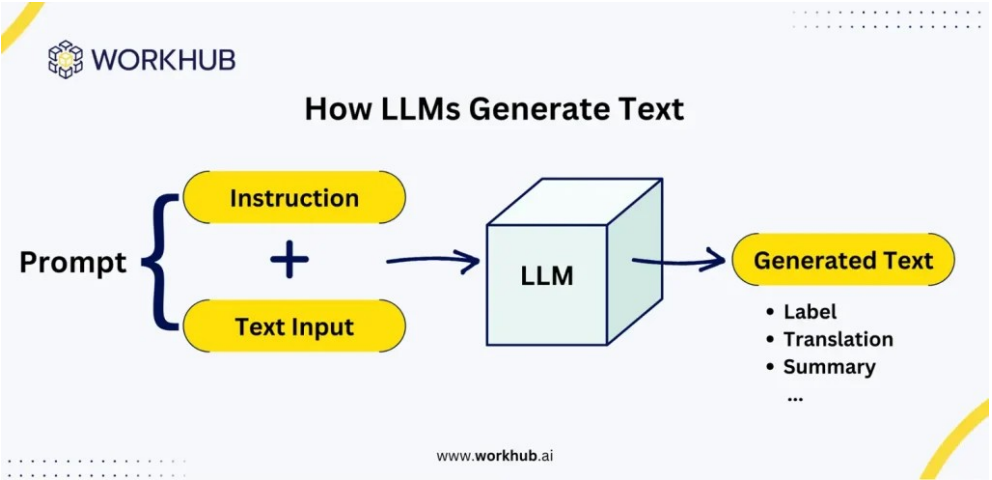


```
Write a Python function:

Function name: validate_email
Input: string email
Output: boolean

Requirements:
- Use regex.
- Reject emails without domain suffix.
- No external libraries.
- Python 3.11 compatible.
```

Code Generation

- Spec-driven prompting
  - LMs respond extremely well to structured specifications
    - Instead of “Build a REST API”
  - Use: the prompt shown
- Why this works.
  - The model has seen
    - API documentation
    - RFC-style specs
    - Technical documentation
    - GitHub READMEs
  - Structured prompts activate those learned patterns

```
Design a REST API with the following specification:

Language: Node.js
Framework: Express
Database: PostgreSQL

Endpoints:
1. POST /users
2. GET /users/{id}
3. DELETE /users/{id}

Constraints:
- Use async/await
- Input validation required
- Return JSON responses
- Proper HTTP status codes
```

Code Generation

- Unit test generation
  - Generating tests from spec
- The model can produce
  - Happy path tests
  - Edge case tests
  - Exception tests

```
Given the following function specification, generate comprehensive pytest unit tests.

Specification:
Function: calculate_discount(price: float, percentage: float) -> float
Requirements:
- percentage must be between 0 and 100
- price must be non-negative
- Raise ValueError for invalid input
```

Code Generation

- Unit test generation
  - Generating tests from code
  - Very useful for
    - Legacy systems
    - Coverage improvement
    - Regression testing
  - After test generation
    - Run tests
    - Detect failures
    - Feed failure back into model
  - This creates an iterative refinement loop

```
Generate unit tests for the following function.
Cover edge cases and failure modes.

<code block>
```

Code Generation

- Refactoring prompts
  - Refactoring is not rewriting
  - It is changing the structure of the code without changing its functionality
  - Weak prompt
    - Improve this code
    - Produces vague results
  - Better prompt shown on the right

```
Refactor the following Python function:

Goals:
- Improve readability
- Follow PEP8
- Reduce cyclomatic complexity
- Preserve behavior exactly
- Add type hints

Do not change functionality.
```

Code Generation

- Advanced refactoring
  - You can request
    - Convert to functional style
    - Extract helper methods
    - Apply design patterns
    - Optimize performance
    - Remove duplication
  - Structured goal-driven refactoring works best

```
Refactor the following Python function:

Goals:
- Improve readability
- Follow PEP8
- Reduce cyclomatic complexity
- Preserve behavior exactly
- Add type hints

Do not change functionality.
```

Code Generation

- Basic debugging
  - Debugging prompts require structured reasoning
  - Providing
    - Expected behavior
    - Actual behavior
    - Reproducible example
  - Significantly improves results

```
Follow this debugging process:
1. Identify potential error sources
2. Explain root cause
3. Provide corrected implementation
4. Suggest test case to prevent regression
```

Code Generation

- Secure coding prompts
  - LLMs default to functionality over security unless instructed otherwise
  - Common security prompt additions
    - “Follow OWASP guidelines”
    - “Prevent XSS”
    - “Use constant-time comparison”
    - “Avoid hardcoded secrets”
    - “Use prepared statements”
- Security must be explicitly requested

```
Write a login handler in Flask.

Security requirements:
- Use parameterized SQL queries
- Hash passwords with bcrypt
- Prevent SQL injection
- Validate input
- Do not expose stack traces
- Follow OWASP best practices
```

Code Generation

- Secure coding prompts
  - Recent Claude models tend to apply common safety practices by default
    - Parameterized queries
    - Input validation
    - No hard-coded secrets
  - For safety-critical code, do not rely on defaults
    - State security requirements explicitly
  - For high-assurance code, use prompting with:
    - Static analysis tools
    - Security-focused code review
    - Explicit threat modeling in the prompt

```
Write a login handler in Flask.

Security requirements:
- Use parameterized SQL queries
- Hash passwords with bcrypt
- Prevent SQL injection
- Validate input
- Do not expose stack traces
- Follow OWASP best practices
```

Code Generation

- Best practice AI assisted pipeline
  - Stages
  - Spec-driven code generation
  - Unit test generation
  - Run tests
  - Debug failures
  - Refactor for readability
  - Add security hardening
- Leverages LLM coding assistant at each step

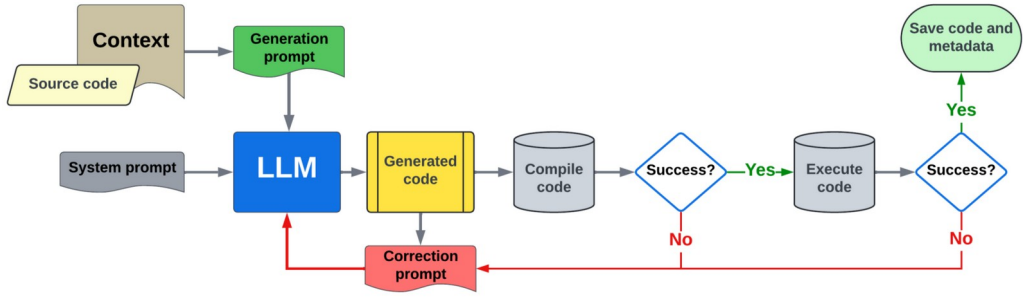

Image Credit: https://www.themoonlight.io/zh/review/lassi-an-llm-based-automated-self-correcting-pipeline-for-translating-parallel-scientific-codes

Code Generation

- This pipeline can be:
  - Implemented manually using one prompt per step
  - Implemented iteratively in a single chat
  - Automated using an agentic coding tool
  - GitHub Copilot agentic surface

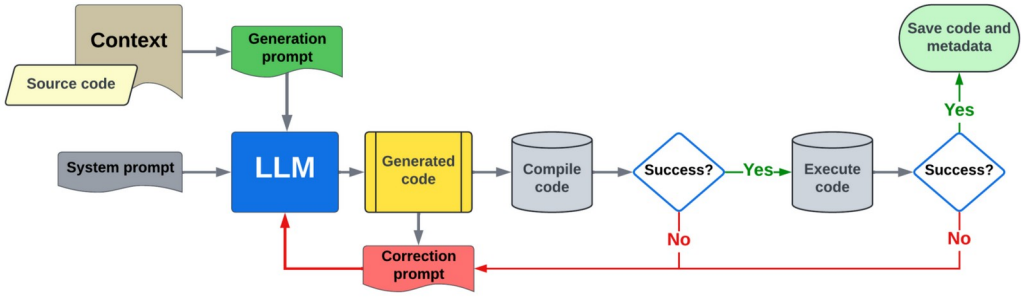

Image Credit: https://www.themoonlight.io/zh/review/lassi-an-llm-based-automated-self-correcting-pipeline-for-translating-parallel-scientific-codes

Prompt Templates

- In a production system
  - Prompts are development artifacts
  - Treated like code artifacts or test artifacts
- They should be
  - Modular
  - Reusable
  - Versioned
  - Testable

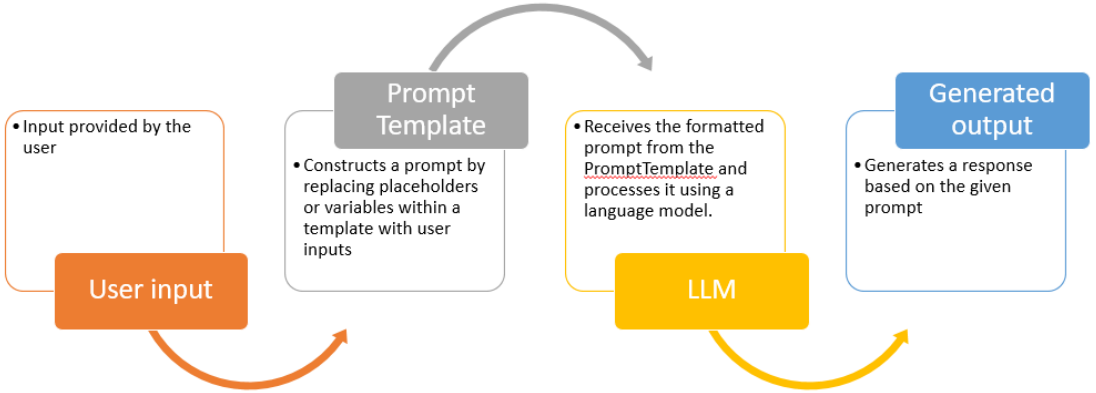

Image Credit:https://vitalflux.com/llm-chain-openai-chatgpt-python-example/

Prompt Templates

- Without templates
  - Prompts become inconsistent
  - Behavior drifts
  - Debugging becomes difficult
  - Updates cause regressions
  - Security risks increase
- With templates
  - Outputs are predictable
  - Prompts are reusable
  - Versioning becomes possible
  - Testing is easier
  - Scaling is manageable
  - Consistent with software engineering

```
system = """You are a {role}.

Follow these constraints:
{constraints}
"""

user = """<task>
{task_description}
</task>

<input>
{input_data}
</input>

<output_format>
{output_format}
</output_format>
"""
```

Prompt Templates

- Sample classification template
- These templates are still valid when the selected model in Copilot is Claude
- In Copilot these templates would typically be stored as:
  - .github/prompts/classify.prompt.md
  - .github/prompts/summarize.prompt.md
  - .github/prompts/generate-function.prompt.md
  - And referenced from chat with /classify, /summarize, etc.

```
System: You are a classification assistant.

User:
<task>
Classify the following text into one of these categories
- Billing
- Technical Support
- Sales
- Other
</task>

<text>
{{customer_message}}
</text>

<output>
Return only the category name. No explanation.
</output>
```

Prompt Templates

- Summarization template

```
System: You are an executive summarization assistant.

User:
<task>
Summarize the following document in {{num_sentences}} sentences.
</task>

<document>
{{document_text}}
</document>

<constraints>
- Focus on key business insights
- Avoid minor details
- Do not add information
</constraints>

<output>
Plain text summary only.
</output>
```

Prompt Templates

- Code generation template

```
System: You are a senior Python engineer.

User:
<task>
Write a function that satisfies the following specification.
</task>

<specification>
Function Name: {{function_name}}
Input: {{input_description}}
Output: {{output_description}}
</specification>

<constraints>
- Python 3.11
- No external libraries
- Include type hints
- Handle edge cases
</constraints>

<output>
Return only the code. No prose.
</output>
```

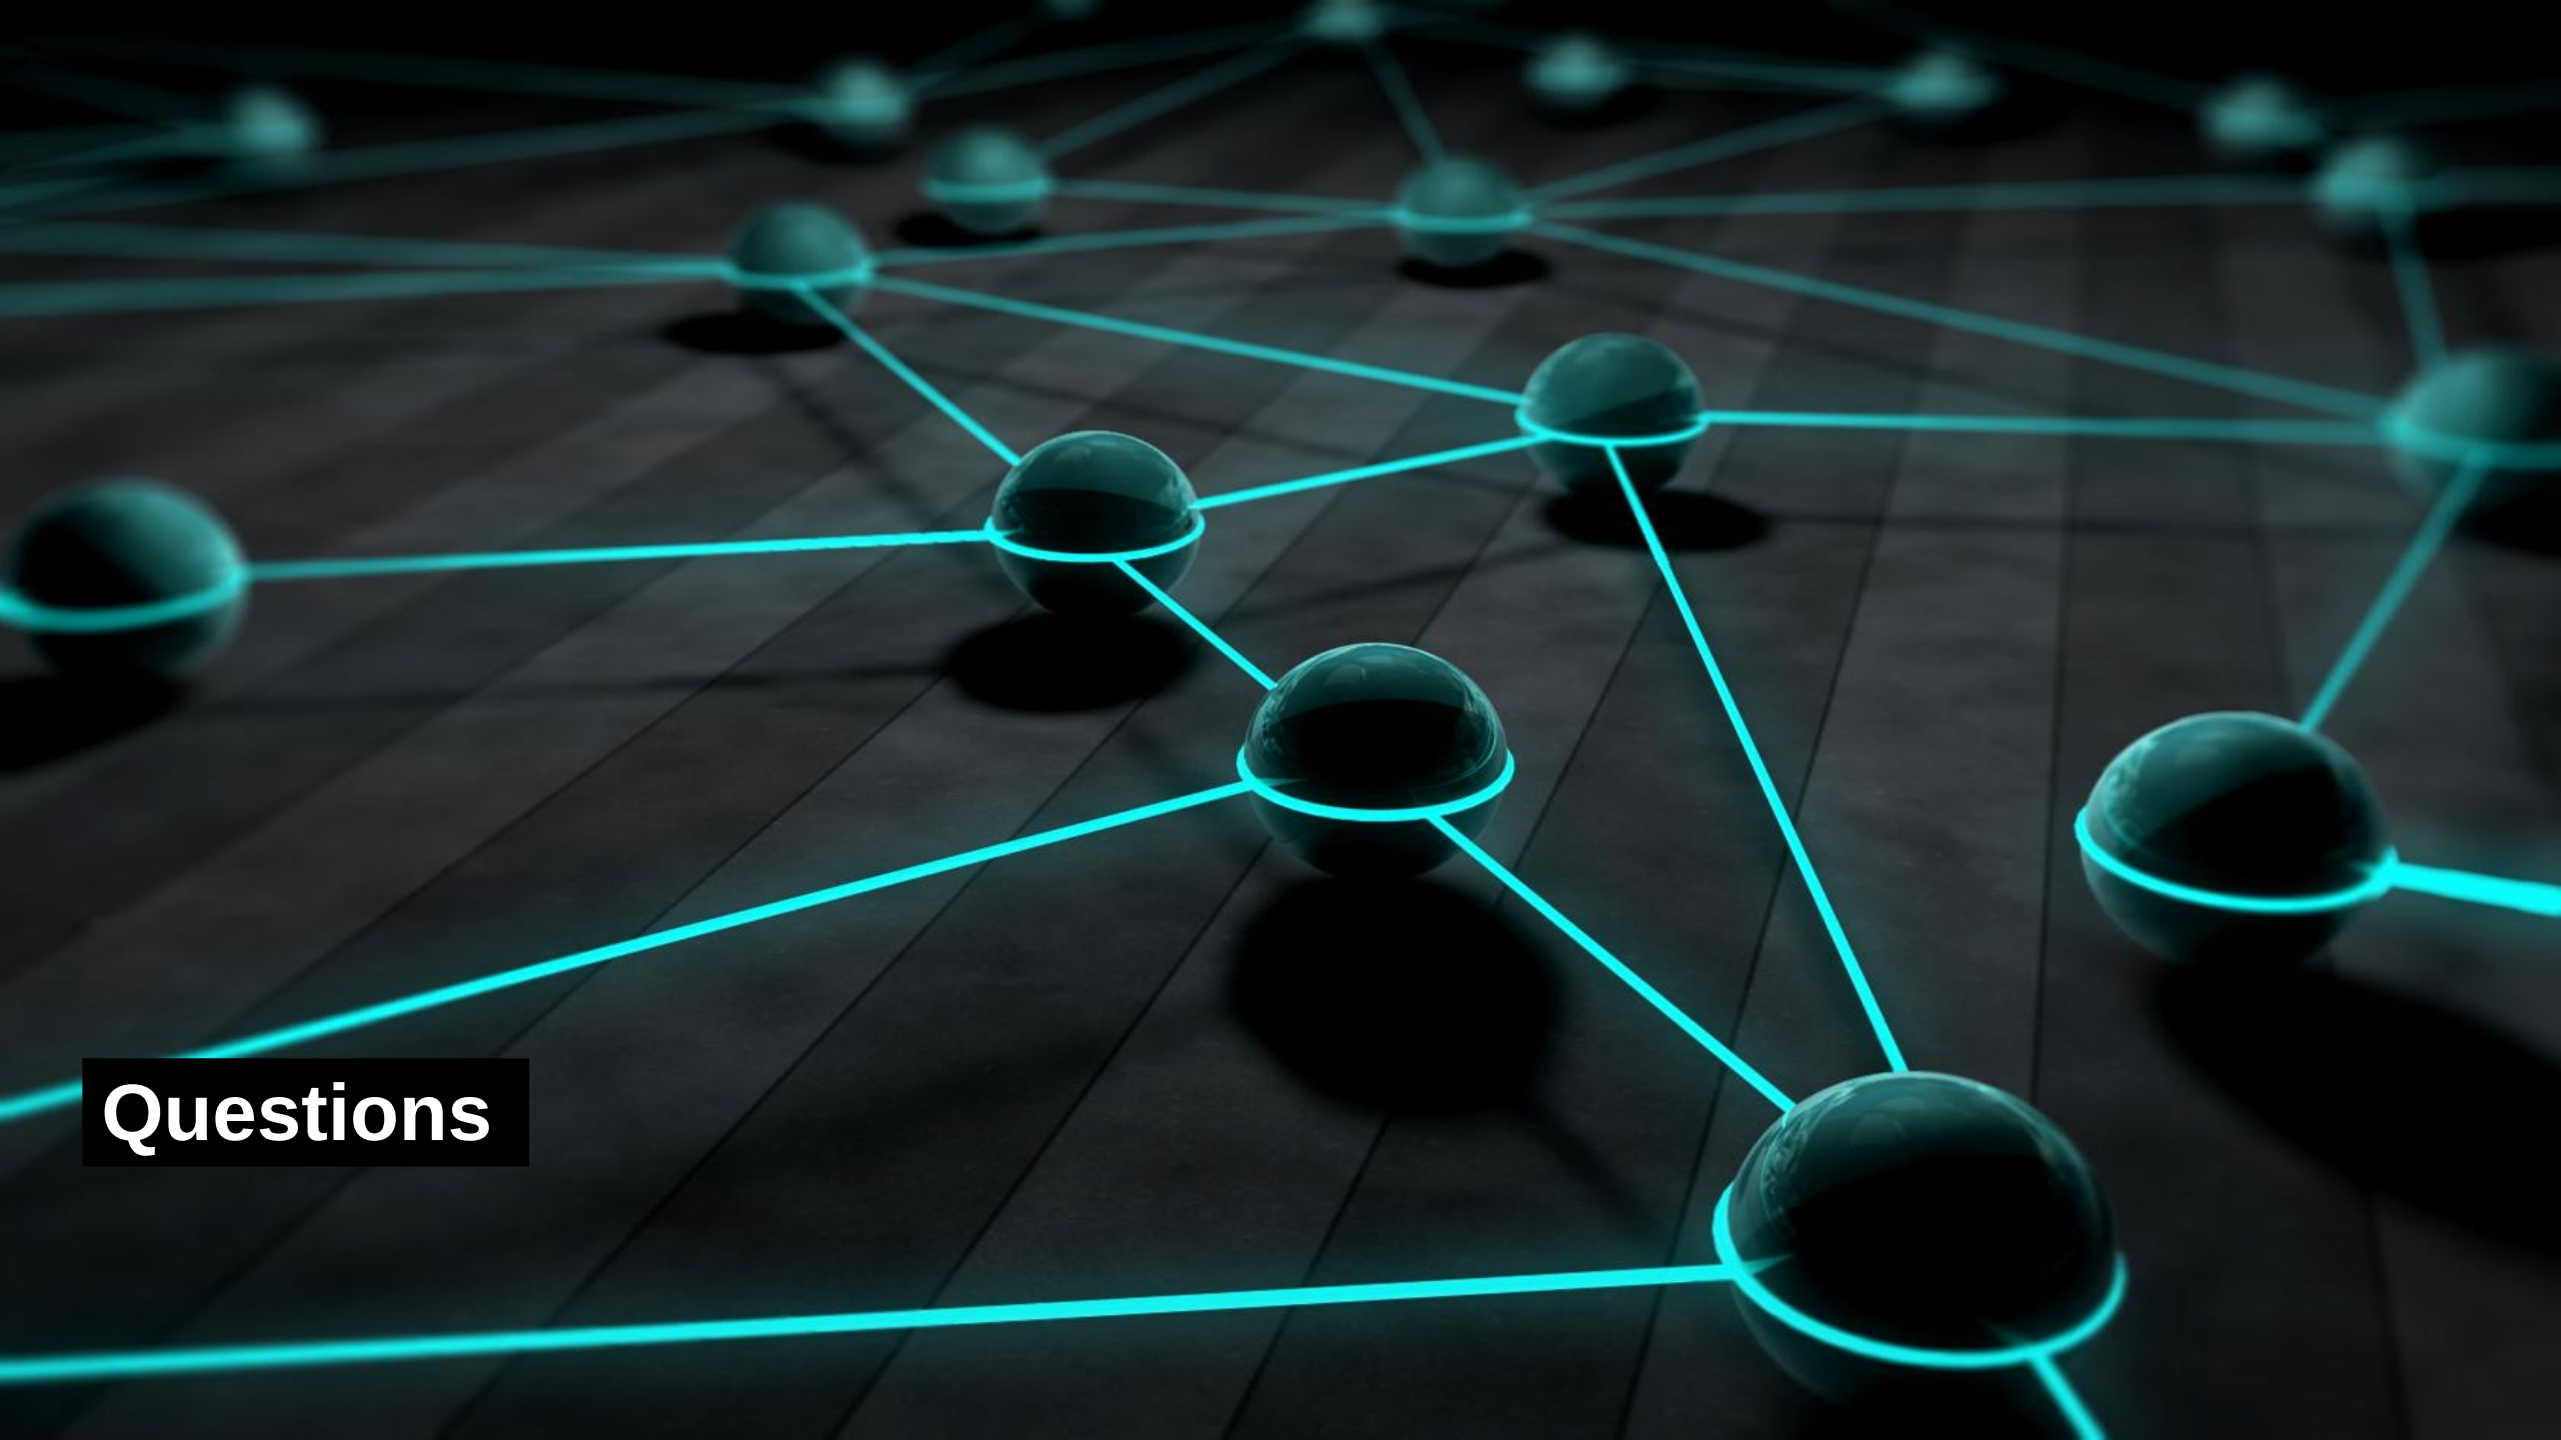
Questions